

# Complete Backend & System Design Interview Guide

Your complete roadmap to crushing backend and system design interviews

---

## Part 1: Backend Fundamentals

### 1. How the Internet Works

**Q: What happens when you type google.com in your browser?**

**Answer:**

1. **DNS Resolution:** Browser finds IP address (142.250.190.46)
  2. **TCP Handshake:** 3-way handshake (SYN, SYN-ACK, ACK)
  3. **TLS Handshake:** Negotiate encryption, verify certificate
  4. **HTTP Request:** Send GET request
  5. **Server Processing:** Load balancer routes to server
  6. **HTTP Response:** Server sends HTML back
  7. **Rendering:** Browser parses HTML, loads assets, renders page
- 

### 2. HTTP Methods

**Q: When to use each HTTP method?**

**Answer:**

- **GET:** Retrieve data. Idempotent, cacheable, no body
  - **POST:** Create new resource. Not idempotent
  - **PUT:** Replace entire resource. Idempotent
  - **PATCH:** Partial update
  - **DELETE:** Remove resource. Idempotent
- 

### 3. HTTP Status Codes

**Answer:**

- **2xx:** Success (200 OK, 201 Created, 204 No Content)
  - **3xx:** Redirect (301 Permanent, 302 Temporary, 304 Not Modified)
  - **4xx:** Client Error (400 Bad Request, 401 Unauthorized, 403 Forbidden, 404 Not Found, 429 Rate Limit)
  - **5xx:** Server Error (500 Internal Error, 502 Bad Gateway, 503 Unavailable, 504 Timeout)
- 

### 4. REST API Design

**Q: Design a blog API**

**Answer:**



GET /posts - List posts  
GET /posts/{id} - Get post  
POST /posts - Create post  
PUT /posts/{id} - Update post  
DELETE /posts/{id} - Delete post  
GET /posts/{id}/comments - Get comments  
POST /posts/{id}/comments - Add comment

Best practices: Use nouns, plurals, proper HTTP methods, versioning (/v1/)

## 5. Authentication vs Authorization

Answer:

- **Authentication:** Who are you? (Login, JWT, OAuth)
- **Authorization:** What can you do? (RBAC, permissions)

Example: User logs in (auth) → tries to delete post (authz checks if they own it)

## 6. Authentication Methods

Session-Based:



python

*# Server stores session, returns session\_id in cookie*  
login() → session\_id → cookie  
request() → cookie → validate session

Pros: Can revoke instantly Cons: Stateful, doesn't scale easily

Token-Based (JWT):



python

# Server generates signed token

login() → JWT token → client stores

request() → JWT in header → validate signature

Pros: Stateless, scales well Cons: Can't revoke until expiry

---

## 7. SQL vs NoSQL

Use SQL when:

- Need ACID transactions (banking)
- Complex queries with JOINS
- Clear relationships
- Schema is stable

Use NoSQL when:

- Massive scale needed
- Schema changes frequently
- Unstructured data
- Eventual consistency OK

NoSQL Types:

- **Document:** MongoDB (JSON docs)
  - **Key-Value:** Redis (cache)
  - **Column-Family:** Cassandra (time series)
  - **Graph:** Neo4j (social networks)
- 

## 8. Database Indexing

**Q: When to use indexes?**

**Answer:** Index = Table of contents for database



sql

```
CREATE INDEX idx_email ON users(email);
```

```
-- Fast: WHERE email = 'john@example.com'
```

```
CREATE INDEX idx_name_age ON users(last_name, first_name);
```

```
-- Fast: WHERE last_name = 'Smith'
```

```
-- Slow: WHERE first_name = 'John' (skips first column)
```

**Trade-offs:**

- Speeds up reads
  - Slows down writes
  - Uses storage
- 

## 9. ACID Properties

**Answer:**

- **Atomicity:** All or nothing
- **Consistency:** Valid state always
- **Isolation:** Concurrent transactions don't interfere
- **Durability:** Committed data survives crashes

**Isolation Levels** (weak → strong):

1. Read Uncommitted (dirty reads possible)
  2. Read Committed (most common)
  3. Repeatable Read
  4. Serializable (slowest, safest)
- 

## 10. N+1 Query Problem

**Problem:**



python

```
posts = db.query("SELECT * FROM posts") # 1 query
for post in posts:
    author = db.query("SELECT * FROM users WHERE id = ?", post.author_id) # N queries
# Total: 1 + N queries!
```

**Solution:**



python

```
# Use JOIN
posts = db.query("""
    SELECT posts.*, users.name
    FROM posts JOIN users ON posts.author_id = users.id
""") # 1 query only
```

---

# 11. Caching Strategies

## Cache-Aside:



python

```
data = cache.get(key)
if not data:
    data = db.get(key)
    cache.set(key, data)
return data
```

## Write-Through:



python

```
db.update(key, data)
cache.set(key, data) # Update cache immediately
```

## Cache Eviction:

- **LRU**: Remove least recently used
- **LFU**: Remove least frequently used
- **TTL**: Expire after time

---

# 12. Message Queues

## Why use them:

- Decouple services
- Handle traffic spikes
- Async processing
- Retry failures

## Example:



python

*# Producer*

```
queue.publish('send_email', {'to': 'user@example.com'})
```

*# Consumer*

```
@worker('send_email')
```

```
def send_email(message):  
    email.send(message['to'])
```

---

## 13. API Rate Limiting

**Token Bucket:**



python

```
class RateLimiter:  
    def allow_request(self):  
        # Refill tokens over time  
        # Check if tokens available  
        if tokens >= 1:  
            tokens -= 1  
            return True  
        return False
```

Return headers:



```
X-RateLimit-Limit: 100
```

```
X-RateLimit-Remaining: 87
```

---

## Part 2: System Design Principles

### 1. Scalability

**Vertical Scaling (Scale Up):**

- Add more CPU/RAM to one machine
- Simple but has limits
- Single point of failure

**Horizontal Scaling (Scale Out):**

- Add more machines
  - Needs load balancing
  - Nearly unlimited scaling
- 

**2. CAP Theorem**

You can only have 2 of 3:

- **Consistency:** All nodes see same data
- **Availability:** Always get response
- **Partition Tolerance:** Works despite network failures

**In practice:**

- **CP:** MongoDB (sacrifice availability)
  - **AP:** Cassandra (sacrifice consistency)
- 

**3. Load Balancing**

**Algorithms:**

- **Round Robin:** Request 1 → A, 2 → B, 3 → C, 4 → A
- **Least Connections:** Route to server with fewest connections
- **IP Hash:** Same client → same server (session affinity)

**Layer 4 vs Layer 7:**

- **L4:** Routes by IP/port (fast)
  - **L7:** Routes by URL/headers (flexible)
- 

**4. Caching Layers**



User → Browser Cache → CDN → App Cache (Redis) → DB Cache → Database

Each layer stores subset of data, getting faster but smaller.

---

**5. Database Sharding**

Split data across multiple databases.

**Hash-Based:**



python

```
shard_id = hash(user_id) % num_shards
```

Even distribution

**Range-Based:**



- Shard 1: user\_id 1-1M
- Shard 2: user\_id 1M-2M

Can have hotspots

---

**6. Replication**

**Master-Slave:**



Master (writes) → Slave1, Slave2, Slave3 (reads)

Scales reads, single point for writes

**Master-Master:**



Master1 ↔ Master2 (both accept writes)

No single point of failure, needs conflict resolution

---

**7. Microservices vs Monolith**

**Monolith:**

- Single codebase/deployment
- Good for: Small teams, early stage
- Scales as one unit

**Microservices:**



- Independent services
- Good for: Large teams, need independent scaling
- Complex infrastructure

Start with modular monolith, split later if needed

---

## 8. Service Communication

Sync (HTTP):



python

```
payment = http.post("payment-service/charge", data)
```

Simple, but tight coupling

Async (Message Queue):



python

```
queue.publish("order.created", data)  
# Payment service subscribes and processes
```

Loose coupling, eventual consistency

---

## 9. Circuit Breaker

Prevents cascading failures.

States:

- **Closed:** Normal (requests pass)
- **Open:** Service failing (fail fast)
- **Half-Open:** Testing recovery



python

```
if circuit_breaker.is_open():
    return error("Service unavailable")
try:
    return call_service()
except:
    circuit_breaker.record_failure()
```

---

## 10. Consistency Models

### Strong Consistency:

- After write, all reads see new value immediately
- Slower
- Use for: Banking, inventory

### Eventual Consistency:

- After write, reads eventually see new value
  - Faster
  - Use for: Social media, analytics
- 

# Part 3: System Design Questions

## Design 1: URL Shortener

**Requirements:** Shorten URLs, redirect, 100M URLs/month

### Solution:

API:



POST /shorten → {short\_url}  
GET /{code} → redirect

Generate short code:



python

```
def encode(id):
    # Base62 encoding (0-9, a-z, A-Z)
    chars = "0-9a-zA-Z"
    result = ""
    while id > 0:
        result = chars[id % 62] + result
        id //= 62
    return result # e.g., "dnh9"
```

Architecture:



User → LB → App Servers → Redis (cache) → DB

Read flow:



python

```
url = cache.get(code)
if not url:
    url = db.get(code)
    cache.set(code, url)
return redirect(url)
```

**Scale:** 40 writes/sec, 4000 reads/sec. Cache hot URLs (80/20 rule).

## Design 2: Twitter Feed

**Requirements:** Post tweets, follow users, view timeline

**Schema:**



sql

```
tweets(id, user_id, content, created_at)
follows(follower_id, followee_id)
```

**Timeline Approaches:**

**Pull (read-time):**



python

```
following_ids = get_following(user_id)
tweets = db.query("SELECT * FROM tweets WHERE user_id IN (?)", following_ids)
```

Slow for many followings

**Push (write-time):**



python

```
# When posting
followers = get_followers(user_id)
for follower in followers:
    cache.lpush(f"timeline:{follower}", tweet_id)
```

Fast reads, expensive for celebrities

**Hybrid (Twitter's approach):**

- Push for regular users
- Pull for celebrities
- Merge at read time

---

**Design 3: Instagram**

**Requirements:** Upload photos, view feed

**Upload flow:**



python

*# 1. Save original to S3*

```
s3.upload(photo_id, file)
```

*# 2. Queue processing*

```
queue.publish('resize_photo', photo_id)
```

*# 3. Worker resizes (thumbnail, medium, large)*

*# 4. Upload all sizes to S3/CDN*

**Feed:** Same as Twitter (pre-compute timelines)

**Optimization:**

- Multiple image sizes
  - CDN for global distribution
  - Lazy loading
- 

## Design 4: Uber

**Requirements:** Match riders with drivers, real-time location

**Key component:** Find nearby drivers

Using Redis Geo:



python

*# Driver updates location*

```
redis.geoadd('drivers', lng, lat, driver_id)
```

*# Find nearby*

```
drivers = redis.georadius('drivers', lng, lat, 5km)
```

**Ride flow:**



1. Rider requests → find nearby drivers
2. Notify drivers
3. Driver accepts
4. Real-time tracking via WebSocket
5. Complete ride → calculate fare

**Fare calculation:**



python

```
fare = base + (distance * per_km) + (time * per_min)
fare *= surge_multiplier
```

**Design 5: Netflix**

**Requirements:** Stream videos, multiple quality levels

**Video processing:**



python

```
# Upload → transcode to 360p, 720p, 1080p, 4K
# Create HLS segments (10 sec chunks)
# Store on S3 → serve via CDN
```

**Adaptive streaming (HLS):**

- Client downloads manifest
- Picks quality based on bandwidth
- Switches quality dynamically

**Architecture:**



```
User → CDN (edge) → Origin → S3
```

**Scale:** Cache video segments (immutable), use geo-distributed CDN

## Design 6: WhatsApp

**Requirements:** Real-time messaging, delivery guarantee

**Architecture:**



User → WebSocket → Chat Server → Kafka → Cassandra

**Message flow:**



python

```
# Send
kafka.publish(recipient_id, message)

# Receive
@websocket.on_message
def handle(msg):
    # If recipient online
    websocket.send(recipient_id, msg)
    # Else store for later
    db.save(msg)
```

**Delivery guarantee:**

- Message ID for deduplication
- ACK from recipient
- Retry with exponential backoff

---

## Design 7: Google Drive

**Requirements:** Upload files, share, sync

**Upload:**



python

*# Split file into chunks*

```
chunks = split_file(file, chunk_size=4MB)
```

*# Upload in parallel*

```
for chunk in chunks:
```

```
    s3.upload(chunk_id, chunk)
```

*# Save metadata*

```
db.save(file_id, chunk_ids)
```

## Sync:

- Client polls for changes (or WebSocket)
- Download only changed chunks
- Merge locally

## Sharing:



sql

```
CREATE TABLE permissions (  
    file_id, user_id, access_level  
)
```

---

## Design 8: Notification System

**Requirements:** Send email, SMS, push notifications

**Architecture:**



API → Queue → Workers (Email/SMS/Push) → Provider APIs

**Priority queue:**



python



*# High priority (OTP)*

```
queue.publish('notifications_high', msg)
```

*# Low priority (marketing)*

```
queue.publish('notifications_low', msg)
```

**Retry logic:**



python

```
@worker
```

```
def send(msg):  
    try:  
        provider.send(msg)  
    except:  
        if retry_count < 3:  
            queue.publish_delayed(msg, delay=2^retry_count)
```

---

**Design 9: Rate Limiter**

**Requirements:** Limit API requests per user

**Token Bucket (Redis):**



python

```
key = f"rate:{user_id}"  
count = redis.incr(key)  
if count == 1:  
    redis.expire(key, window)  
if count > limit:  
    return 429 # Too many requests
```

**Distributed:** Use Redis cluster, handle race conditions with Lua scripts

---

**Design 10: Search Engine**

**Requirements:** Index web pages, search fast

Crawling:



python

```
queue = ['seed_urls']
while queue:
    url = queue.pop()
    page = fetch(url)
    index(page)
    queue.extend(extract_links(page))
```

Indexing (Inverted Index):



"hello" → [doc1, doc5, doc9]  
"world" → [doc1, doc3, doc9]

Ranking:



python

score = TF-IDF \* PageRank \* freshness

Architecture:



User → LB → Query Servers → Index Shards

Part 4: Key Concepts Summary

Performance Metrics

- **Latency:** Time to complete one request
- **Throughput:** Requests per second
- **Availability:** Uptime percentage (99.9% = 43 min downtime/month)

# Back-of-Envelope Calculations



1 million requests/day = 12 requests/sec

1 billion requests/day = 12K requests/sec

1 KB \* 1 million users = 1 GB

1 MB \* 1 million users = 1 TB

## Storage



Text: ~1 KB per message

Image: ~500 KB

Video: ~50 MB (compressed)

## Common Numbers



L1 cache: 1 ns

RAM: 100 ns

SSD: 100 μs

Network within DC: 500 μs

HDD: 10 ms

Network cross-continent: 100 ms

---

# Part 5: Interview Tips

## How to Approach System Design

1. **Clarify requirements** (5 min)
  - Functional: What features?
  - Non-functional: Scale, latency, availability?
  - Constraints: Budget, time?
2. **High-level design** (10 min)
  - Draw basic components
  - API design

- Database schema
- 3. **Deep dive** (20 min)
  - Focus on 2-3 components
  - Discuss trade-offs
  - Address bottlenecks
- 4. **Wrap up** (5 min)
  - Monitoring/alerts
  - Future improvements

Things to Mention

- **Bottlenecks:** Where will system fail at scale?
- **Trade-offs:** Why choose X over Y?
- **Monitoring:** How to detect issues?
- **Failure scenarios:** What if DB goes down?

Red Flags to Avoid

- Don't jump to solution immediately
- Don't over-engineer for day 1
- Don't ignore scale numbers
- Don't forget about failures

---

Part 6: Common Patterns

API Design Patterns

Pagination:



```
GET /posts?page=2&limit=20
GET /posts?cursor=abc123&limit=20 (preferred)
```

Filtering:



```
GET /posts?status=published&author=john
```

Sorting:



```
GET /posts?sort=created_at&order=desc
```

## Versioning:



```
/v1/posts
```

```
Accept: application/vnd.api.v2+json
```

## Database Patterns

### Soft Delete:



```
sql
```

```
UPDATE users SET deleted_at = NOW() WHERE id = ?
```

### Optimistic Locking:



```
sql
```

```
UPDATE orders SET status = 'shipped', version = version + 1  
WHERE id = ? AND version = ?
```

**Denormalization:** Store computed values to avoid JOINS



```
sql
```

```
posts(id, title, author_name) -- Duplicate author name
```

## Caching Patterns

**Cache warming:** Pre-populate cache before launch

**Cache stampede prevention:**



python

```
lock = redis.lock(key)
if lock.acquire():
    data = compute_expensive()
    cache.set(key, data)
    lock.release()
```

## Part 7: Technology Choices

### Databases

**PostgreSQL:** Strong ACID, complex queries **MySQL:** Proven, good for reads **MongoDB:** Flexible schema, document storage **Cassandra:** Write-heavy, time series **Redis:** In-memory cache, pub/sub **Elasticsearch:** Full-text search, analytics

### Message Queues

**RabbitMQ:** Traditional, reliable **Kafka:** High throughput, event streaming **SQS:** Managed, AWS **Redis Pub/Sub:** Simple, fast

### Caching

**Redis:** Most popular, rich features **Memcached:** Simple, fast **CDN:** CloudFlare, CloudFront

### Storage

**S3:** Object storage, cheap **EBS:** Block storage for servers **Glacier:** Archive, very cheap

## Part 8: Final Checklist

Before your interview, make sure you can:

#### Explain clearly:

- ☐ HTTP request lifecycle
- ☐ Database indexes
- ☐ Caching strategies
- ☐ CAP theorem
- ☐ Sharding vs replication
- ☐ Microservices trade-offs
- ☐ Message queues

#### Design from scratch:

- ☐ URL shortener

- ☐ Social media feed
- ☐ Chat application
- ☐ Video streaming
- ☐ Ride sharing

### Calculate:

- ☐ Storage needs
- ☐ Bandwidth needs
- ☐ QPS (queries per second)
- ☐ Number of servers needed

### Discuss trade-offs:

- ☐ SQL vs NoSQL
- ☐ Sync vs Async
- ☐ Consistency vs Availability
- ☐ Monolith vs Microservices

---

## You're Ready!

You now have comprehensive knowledge of: ✓ Backend fundamentals ✓ System design principles  
✓ Real-world design patterns ✓ Interview strategies

**Practice:** Draw architectures on paper, explain out loud, time yourself.

**Remember:** There's no single "correct" answer. Show your thought process, discuss trade-offs, and be ready to defend your choices.

**Go crush those interviews!** 🚀