

Complete Docker Mastery Guide - Interview Ready

Table of Contents

1. [Docker Fundamentals](#)
 2. [Core Concepts](#)
 3. [Docker Images](#)
 4. [Docker Containers](#)
 5. [Dockerfile Deep Dive](#)
 6. [Docker Networking](#)
 7. [Docker Volumes & Storage](#)
 8. [Docker Compose](#)
 9. [Multi-Stage Builds](#)
 10. [Docker Security](#)
 11. [Docker in Production](#)
 12. [Advanced Topics](#)
 13. [Troubleshooting](#)
 14. [Interview Questions & Answers](#)
-

Docker Fundamentals

What is Docker?

Docker is a platform that packages your application and all its dependencies into containers. Think of it like shipping containers for code. Just like how physical shipping containers can be moved from ships to trucks to trains without unpacking, Docker containers can run anywhere: your laptop, a server, the cloud.

Why Docker Exists

The Problem Before Docker:

- "It works on my machine" syndrome
- Different environments (dev, staging, production) had different configurations
- Dependencies hell - installing the right versions of libraries, languages, tools
- Wasted resources running full virtual machines


What Docker Solves:

- Consistent environments across all stages
- Faster deployment (seconds vs minutes)
- Better resource utilization than VMs
- Easier scaling and orchestration
- Simplified dependency management


Docker vs Virtual Machines

Virtual Machines:





Docker Containers:



Key Differences:

- Containers share the host OS kernel, VMs include full OS
- Containers start in seconds, VMs take minutes
- Containers use MBs of space, VMs use GBs
- You can run 10x more containers than VMs on the same hardware

Core Concepts

1. Images

A Docker image is a read-only template. It's like a class in programming. Contains:

- Base operating system (usually minimal, like Alpine Linux)
- Application code
- Dependencies and libraries
- Configuration files
- Environment variables

2. Containers

A container is a running instance of an image. It's like an object in programming. Features:

- Isolated process with its own filesystem
- Has its own network interface
- Can be started, stopped, moved, deleted
- Ephemeral by default (data lost when deleted)

3. Docker Engine

The core Docker software that:

- Builds images
- Runs containers
- Manages Docker objects (images, containers, networks, volumes)

Components:

- **Docker Daemon** (dockerd): Background service that manages Docker objects
- **Docker CLI** (docker): Command-line tool you use
- **REST API**: How CLI talks to daemon

4. Docker Registry

A storage system for Docker images:

- **Docker Hub**: Public registry (like GitHub for Docker images)
- **Private Registries**: For internal use (AWS ECR, Google GCR, Harbor)

5. Dockerfile

A text file with instructions to build an image. Think of it as a recipe.

Docker Images

Understanding Layers

Docker images are built in layers. Each instruction in a Dockerfile creates a new layer.



dockerfile

```
FROM ubuntu:20.04      # Layer 1: Base OS
RUN apt-get update      # Layer 2: Updated package lists
RUN apt-get install python # Layer 3: Python installed
COPY app.py /app/       # Layer 4: Your code
CMD ["python", "/app/app.py"] # Layer 5: Startup command
```

Why Layers Matter:

- Layers are cached - rebuilds are fast
- Layers are shared between images - saves disk space
- Only changed layers are rebuilt

Layer Caching Strategy:

- Put things that change rarely at the top

- Put things that change often at the bottom
- This keeps cache hits high

Image Commands



bash

List images

`docker images`

`docker image ls`

Pull an image

`docker pull nginx:latest`

`docker pull ubuntu:20.04`

Build an image

`docker build -t myapp:v1 .`

`docker build -t myapp:v1 -f Dockerfile.prod .`

Tag an image

`docker tag myapp:v1 myapp:latest`

`docker tag myapp:v1 myregistry.com/myapp:v1`

Push to registry

`docker push myregistry.com/myapp:v1`

Remove images

`docker rmi nginx:latest`

`docker image rm myapp:v1`

Remove unused images

`docker image prune`

`docker image prune -a` *# Remove all unused images*

Inspect image

`docker image inspect nginx:latest`

View image history (layers)

`docker history nginx:latest`

Save/Load images (for offline transfer)

`docker save myapp:v1 > myapp.tar`

`docker load < myapp.tar`

Export/Import (from container)

```
docker export container_name > container.tar
docker import container.tar myapp:imported
```

Image Naming Convention



[registry]/[username]/[repository]:[tag]

Examples:

- nginx:latest # Official image from Docker Hub
- ubuntu:20.04 # Official Ubuntu image
- myusername/myapp:v1.0 # Your image on Docker Hub
- gcr.io/myproject/myapp:latest # Google Container Registry
- localhost:5000/myapp:dev # Local private registry

Docker Containers

Container Lifecycle



Created → Running → Paused → Stopped → Deleted

Essential Container Commands



bash

Run a container

`docker run nginx`

`docker run -d nginx` *# Detached (background)*

`docker run -d --name web nginx` *# Named container*

`docker run -d -p 8080:80 nginx` *# Port mapping*

`docker run -d -e ENV_VAR=value nginx` *# Environment variable*

`docker run -it ubuntu bash` *# Interactive with terminal*

List containers

`docker ps` *# Running containers*

`docker ps -a` *# All containers (including stopped)*

`docker ps -q` *# Only container IDs*

Start/Stop containers

`docker start container_name`

`docker stop container_name`

`docker restart container_name`

`docker pause container_name` *# Pause processes*

`docker unpause container_name`

Remove containers

`docker rm container_name`

`docker rm -f container_name` *# Force remove running container*

`docker container prune` *# Remove all stopped containers*

Execute commands in running container

`docker exec container_name ls /app`

`docker exec -it container_name bash` *# Interactive shell*

View logs

`docker logs container_name`

`docker logs -f container_name` *# Follow logs (like tail -f)*

`docker logs --tail 100 container_name`

`docker logs --since 10m container_name`

View container details

`docker inspect container_name`

`docker stats` *# Resource usage (live)*

`docker top container_name` *# Running processes*

Copy files

```
docker cp file.txt container_name:/path/
docker cp container_name:/path/file.txt ./
```

View port mappings

```
docker port container_name
```

Attach to running container

```
docker attach container_name
```

Docker Run Options (Deep Dive)



bash

Full example with common options

```
docker run \
  -d \                # Detached mode
  --name myapp \      # Container name
  --restart unless-stopped \ # Restart policy
  -p 8080:80 \        # Port mapping host:container
  -p 127.0.0.1:8081:81 \ # Bind to specific IP
  -e DATABASE_URL=postgres://... \ # Environment variable
  --env-file .env \   # Load env vars from file
  -v /host/path:/container/path \ # Volume mount
  -v myvolume:/data \ # Named volume
  --mount type=bind,src=...,dst=... \ # Alternative mount syntax
  --network mynetwork \ # Custom network
  --network-alias dbserver \ # Network alias
  -w /app \            # Working directory
  --user 1000:1000 \   # Run as specific user
  -m 512m \            # Memory limit
  --cpus 2 \          # CPU limit
  --health-cmd "curl -f http://localhost || exit 1" \
  --health-interval 30s \ # Health check
  --log-driver json-file \ # Logging driver
  --log-opt max-size=10m \ # Log options
  myimage:tag \        # Image
  python app.py         # Command to run
```


Restart Policies



bash

no: Don't restart (default)

`docker run --restart no nginx`

always: Always restart if stopped

`docker run --restart always nginx`

on-failure: Restart if exit code is non-zero

`docker run --restart on-failure nginx`

`docker run --restart on-failure:5 nginx # Max 5 retries`

unless-stopped: Always restart unless manually stopped

`docker run --restart unless-stopped nginx`

Dockerfile Deep Dive

Dockerfile Structure



dockerfile

Syntax version (optional but recommended)

syntax=docker/dockerfile:1

Base image

FROM node:18-alpine

Metadata

LABEL maintainer="you@example.com"

LABEL version="1.0"

LABEL description="My awesome app"

Environment variables

ENV NODE_ENV=production

ENV APP_PORT=3000

Working directory

WORKDIR /app

Copy files

COPY package*.json ./

COPY . .

Run commands (during build)

RUN npm install --production

RUN npm run build

Expose ports (documentation only)

EXPOSE 3000

Create volumes

VOLUME /app/data

Set user (don't run as root)

USER node

Default command (runtime)

CMD ["node", "server.js"]

All Dockerfile Instructions

FROM



dockerfile

```
# Start from base image
FROM ubuntu:20.04

# Multi-stage build
FROM node:18 AS builder
FROM nginx:alpine AS runtime

# Platform-specific
FROM --platform=linux/amd64 ubuntu:20.04
```

RUN



dockerfile

```
# Shell form (runs in /bin/sh -c)
RUN apt-get update && apt-get install -y curl

# Exec form (doesn't invoke shell)
RUN ["/bin/bash", "-c", "echo hello"]

# Multi-line for readability
RUN apt-get update && apt-get install -y \
    curl \
    git \
    vim \
    && rm -rf /var/lib/apt/lists/*
```

COPY vs ADD



dockerfile

COPY - simple file copy (preferred)

COPY package.json /app/

COPY . /app/

ADD - COPY with superpowers (use sparingly)

ADD https://example.com/file.txt /app/ # Can download URLs

ADD archive.tar.gz /app/ # Auto-extracts archives

COPY with ownership

COPY --chown=node:node . /app/

CMD vs ENTRYPOINT

CMD: Default command, can be overridden



dockerfile

CMD ["python", "app.py"]

docker run myimage -> runs python app.py

docker run myimage bash -> runs bash (CMD overridden)

ENTRYPOINT: Command that always runs



dockerfile

ENTRYPOINT ["python", "app.py"]

docker run myimage -> runs python app.py

docker run myimage --debug -> runs python app.py --debug

Combined (best practice):



dockerfile

```
ENTRYPOINT ["python", "app.py"]
```

```
CMD ["--mode", "production"]
```

```
# docker run myimage -> python app.py --mode production
```

```
# docker run myimage --mode dev -> python app.py --mode dev
```

ENV vs ARG

ENV: Available at build AND runtime



dockerfile

```
ENV APP_PORT=3000
```

```
# Can use: $APP_PORT
```

```
# Persists in final image
```

ARG: Only available at build time



dockerfile

```
ARG NODE_VERSION=18
```

```
FROM node:${NODE_VERSION}
```

```
# docker build --build-arg NODE_VERSION=20 .
```

```
# Does NOT persist in final image
```

WORKDIR



dockerfile

```
WORKDIR /app
```

```
# All subsequent commands run in /app
```

```
# Creates directory if doesn't exist
```

```
RUN pwd # prints /app
```

```
COPY . . # copies to /app
```

EXPOSE



dockerfile

```
EXPOSE 3000
EXPOSE 8080/tcp
EXPOSE 8081/udp
# Documentation only - doesn't actually publish ports
# Still need -p flag: docker run -p 3000:3000 myimage
```

VOLUME



dockerfile

```
VOLUME /app/data
# Creates mount point
# Data persists even if container is deleted
```

USER



dockerfile

```
USER node
# All subsequent commands run as 'node' user
# Container also runs as this user
# Security best practice - don't run as root
```

HEALTHCHECK



dockerfile

```
HEALTHCHECK --interval=30s --timeout=3s --start-period=5s --retries=3 \
  CMD curl -f http://localhost/ || exit 1
```

or

```
HEALTHCHECK NONE # Disable healthcheck from base image
```

Best Practices for Dockerfiles

1. Use Specific Tags



dockerfile

Bad

```
FROM node:latest
```

Good

```
FROM node:18.17-alpine
```

2. Minimize Layers



dockerfile

Bad - creates 3 layers

```
RUN apt-get update
```

```
RUN apt-get install -y curl
```

```
RUN apt-get install -y git
```

Good - creates 1 layer

```
RUN apt-get update && apt-get install -y \
  curl \
  git \
  && rm -rf /var/lib/apt/lists/*
```

3. Order Matters (Caching)



dockerfile

```
# Bad - code changes invalidate dependency layer
COPY . /app/
RUN npm install

# Good - dependencies cached unless package.json changes
COPY package*.json /app/
RUN npm install
COPY . /app/
```

4. Use .dockerignore



```
# .dockerignore file
node_modules
.git
.env
*.log
.DS_Store
dist
coverage
```

5. Don't Run as Root



dockerfile

```
# Create non-root user
RUN groupadd -r appuser && useradd -r -g appuser appuser
USER appuser
```

6. Keep Images Small



dockerfile

Use Alpine variants

FROM node:18-alpine

Remove package manager cache

RUN apt-get update && apt-get install -y curl \
&& rm -rf /var/lib/apt/lists/*

Use multi-stage builds (see below)

Real-World Dockerfile Examples

Node.js Application



dockerfile

FROM node:18-alpine

WORKDIR /app

Copy package files

COPY package*.json ./

Install dependencies

RUN npm ci --only=production

Copy application code

COPY . .

Create non-root user

RUN addgroup -g 1001 -S nodejs && \
adduser -S nodejs -u 1001

USER nodejs

EXPOSE 3000

CMD ["node", "server.js"]

Python Application



dockerfile

```
FROM python:3.11-slim

WORKDIR /app

# Install dependencies
COPY requirements.txt .
RUN pip install --no-cache-dir -r requirements.txt

# Copy application
COPY . .

# Create non-root user
RUN useradd -m -u 1000 appuser && \
    chown -R appuser:appuser /app

USER appuser

EXPOSE 8000

CMD ["python", "app.py"]
```

Go Application



dockerfile

FROM golang:1.21-alpine **AS** builder

WORKDIR /app

COPY go.mod go.sum ./

RUN go mod download

COPY . .

RUN CGO_ENABLED=0 GOOS=linux go build -o main .

Final stage

FROM alpine:latest

RUN apk --no-cache add ca-certificates

WORKDIR /root/

COPY --from=builder /app/main .

EXPOSE 8080

CMD ["/main"]

Docker Networking

Network Types

1. Bridge (Default)

- Default network for containers
- Containers can communicate with each other by IP
- Need port publishing to access from host



bash

docker run -d --network bridge nginx

2. Host

- Container shares host's network

- No isolation
- Better performance (no network translation)



bash

```
docker run -d --network host nginx
# nginx accessible on host's port 80 directly
```

3. None

- No networking
- Completely isolated



bash

```
docker run -d --network none nginx
```

4. Custom Bridge (Recommended)

- Better isolation
- Automatic DNS resolution between containers
- Can control which containers can communicate



bash

```
docker network create myapp-network
docker run -d --network myapp-network --name web nginx
docker run -d --network myapp-network --name api node:18
# 'web' can reach 'api' by name: http://api:3000
```

Network Commands



bash

List networks

`docker network ls`

Create network

`docker network create mynetwork`

`docker network create --driver bridge mynetwork`

`docker network create --subnet 172.18.0.0/16 mynetwork`

Inspect network

`docker network inspect mynetwork`

Connect/disconnect container

`docker network connect mynetwork container_name`

`docker network disconnect mynetwork container_name`

Remove network

`docker network rm mynetwork`

`docker network prune` *# Remove unused networks*

Container Communication

Same Network



bash

Create network

```
docker network create app-net
```

Run database

```
docker run -d \  
  --name postgres \  
  --network app-net \  
  -e POSTGRES_PASSWORD=secret \  
  postgres:15
```

Run application (can reach postgres by name)

```
docker run -d \  
  --name api \  
  --network app-net \  
  -e DATABASE_URL=postgres://postgres:secret@postgres:5432/mydb \  
  myapi:latest
```

DNS Resolution

Containers on custom networks can reach each other by:

- Container name
- Network aliases



bash

```
docker run -d \  
  --name db \  
  --network app-net \  
  --network-alias database \  
  --network-alias db-server \  
  postgres:15
```

Can connect using: db, database, or db-server

Port Publishing



bash

Publish single port

```
docker run -p 8080:80 nginx
```

Host:Container (localhost:8080 -> container:80)

Publish to specific IP

```
docker run -p 127.0.0.1:8080:80 nginx
```

Only accessible on localhost

Publish all exposed ports to random host ports

```
docker run -P nginx
```

Multiple ports

```
docker run -p 8080:80 -p 8443:443 nginx
```

UDP ports

```
docker run -p 53:53/udp dns-server
```

Network Troubleshooting



bash

See which network container is on

```
docker inspect container_name | grep -A 10 NetworkSettings
```

Test connectivity

```
docker exec container1 ping container2
```

```
docker exec container1 curl http://container2:8080
```

DNS lookup

```
docker exec container1 nslookup container2
```

View network details

```
docker network inspect bridge
```

Docker Volumes & Storage

Why Volumes?

- Container filesystem is ephemeral (deleted with container)

- Volumes persist data outside container lifecycle
- Can share data between containers
- Better performance than bind mounts on Mac/Windows

Volume Types

1. Named Volumes (Recommended)

Managed by Docker, stored in Docker's storage directory.



bash

Create volume

`docker volume create mydata`

Use volume

```
docker run -d \  
  --name postgres \  
  -v mydata:/var/lib/postgresql/data \  
  postgres:15
```

List volumes

`docker volume ls`

Inspect volume

`docker volume inspect mydata`

Remove volume

`docker volume rm mydata`

`docker volume prune` *# Remove unused volumes*

2. Bind Mounts

Mount host directory into container.



bash

Absolute path required

```
docker run -d \  
-v /host/path:/container/path \  
nginx
```

Current directory

```
docker run -d \  
-v $(pwd):/app \  
node:18
```

Read-only mount

```
docker run -d \  
-v $(pwd):/app:ro \  
nginx
```

3. tmpfs Mounts

Store in host memory (not persisted).



bash

```
docker run -d \  
--tmpfs /tmp:rw,size=100m \  
nginx
```

Mount Syntax Comparison



bash

Old style (-v)

```
docker run -v myvolume:/data myimage
```

New style (--mount) - more explicit, recommended

```
docker run \  
  --mount type=volume,source=myvolume,target=/data \  
  myimage
```

```
docker run \  
  --mount type=bind,source=$(pwd),target=/app \  
  myimage
```

```
docker run \  
  --mount type=tmpfs,target=/tmp,tmpfs-size=100m \  
  myimage
```

Volume Commands



bash

Create with options

```
docker volume create \  
  --driver local \  
  --opt type=none \  
  --opt device=/host/path \  
  --opt o=bind \  
  myvolume
```

Backup volume

```
docker run --rm \  
  -v myvolume:/data \  
  -v $(pwd):/backup \  
  ubuntu \  
  tar czf /backup/mydata.tar.gz /data
```

Restore volume

```
docker run --rm \  
  -v myvolume:/data \  
  -v $(pwd):/backup \  
  ubuntu \  
  tar xzf /backup/mydata.tar.gz -C /
```

Copy data between volumes

```
docker run --rm \  
  -v oldvolume:/from \  
  -v newvolume:/to \  
  alpine \  
  sh -c "cp -av /from/* /to/"
```

Storage Drivers

Different storage drivers for the container's writable layer:

- overlay2: Default and recommended
- aufs: Older systems
- btrfs: For btrfs filesystems
- devicemapper: For production on older kernels
- zfs: For ZFS filesystems



bash

Check storage driver

`docker info | grep "Storage Driver"`

Best Practices



bash

Use named volumes for databases

```
docker run -d \  
-v postgres_data:/var/lib/postgresql/data \  
postgres:15
```

Use bind mounts for development

```
docker run -d \  
-v $(pwd):/app \  
node:18
```

Never store data in container writable layer

Always use volumes

Clean up unused volumes regularly

```
docker volume prune
```

Docker Compose

Docker Compose lets you define multi-container applications in a YAML file.

Why Docker Compose?

- Define entire stack in one file
- One command to start everything
- Easier than remembering long docker run commands
- Version control your infrastructure

Basic docker-compose.yml



yaml

version: '3.8'

services:

web:

image: nginx:alpine

ports:

- "8080:80"

volumes:

- ./html:/usr/share/nginx/html

depends_on:

- api

networks:

- frontend

api:

build: ./api

environment:

- DATABASE_URL=postgres://user:pass@db:5432/mydb

depends_on:

- db

networks:

- frontend

- backend

db:

image: postgres:15

environment:

- POSTGRES_PASSWORD=secret

- POSTGRES_DB=mydb

volumes:

- db_data:/var/lib/postgresql/data

networks:

- backend

volumes:

db_data:

networks:

frontend:
backend:

Compose File Structure

Services



yaml

services:

myapp:

Use existing image

image: nginx:alpine

Or build from Dockerfile

build: .

Or build with context and Dockerfile name

build:

context: ./api

dockerfile: Dockerfile.prod

args:

- NODE_ENV=production

Container name

container_name: my-app

Restart policy

restart: unless-stopped

Ports

ports:

- "8080:80" *# host:container*

- "127.0.0.1:8081:81" *# bind to localhost*

Environment variables

environment:

- NODE_ENV=production

- API_KEY=secret

Or from file

env_file:

- .env

- .env.prod

Volumes

volumes:

- ./code:/app *# bind mount*

- app_data:/data *# named volume*

- /etc/config:/etc/config:ro *# read-only*

Networks

networks:

- frontend
- backend

Dependencies

depends_on:

- db
- redis

Command override

command: npm start

Entrypoint override

entrypoint: /app/entrypoint.sh

Working directory

working_dir: /app

User

user: "1000:1000"

Health check

healthcheck:

test: ["CMD", "curl", "-f", "http://localhost"]

interval: 30s

timeout: 10s

retries: 3

start_period: 40s

Resource limits

deploy:

resources:

limits:

cpus: '0.5'

memory: 512M

reservations:

cpus: '0.25'

memory: 256M

Networks



yaml

```
networks:
  frontend:
    driver: bridge

  backend:
    driver: bridge
    internal: true # No external access

  custom:
    driver: bridge
    ipam:
      config:
        - subnet: 172.20.0.0/16
```

Volumes



yaml

```
volumes:
  db_data:
    driver: local

  nfs_data:
    driver: local
    driver_opts:
      type: nfs
      o: addr=10.0.0.1,rw
      device: ":/path/to/dir"
```

Compose Commands



bash

Start services

`docker-compose up`

`docker-compose up -d` *# Detached*

`docker-compose up --build` *# Rebuild images*

`docker-compose up --force-recreate` *# Recreate containers*

Stop services

`docker-compose stop`

`docker-compose down` *# Stop and remove*

`docker-compose down -v` *# Stop and remove volumes*

`docker-compose down --rmi all` *# Stop and remove images*

View logs

`docker-compose logs`

`docker-compose logs -f` *# Follow*

`docker-compose logs -f api` *# Specific service*

List services

`docker-compose ps`

Execute command

`docker-compose exec api bash`

`docker-compose exec db psql -U postgres`

Run one-off command

`docker-compose run api npm test`

`docker-compose run --rm api npm test` *# Remove after*

Scale services

`docker-compose up -d --scale api=3`

Validate compose file

`docker-compose config`

Build images

`docker-compose build`

`docker-compose build --no-cache api`

Pull images

`docker-compose pull`

```
# Restart services
docker-compose restart
docker-compose restart api
```

Environment Variables

.env File



bash

```
# .env
POSTGRES_VERSION=15
API_PORT=3000
NODE_ENV=production
```

docker-compose.yml



yaml

```
services:
  db:
    image: postgres:${POSTGRES_VERSION}

  api:
    build: .
    ports:
      - "${API_PORT}:3000"
    environment:
      - NODE_ENV=${NODE_ENV}
```

Real-World Examples

Full Stack Application



yaml

version: '3.8'

services:

nginx:

image: nginx:alpine

ports:

- "80:80"
- "443:443"

volumes:

- ./nginx.conf:/etc/nginx/nginx.conf:ro
- ./ssl:/etc/nginx/ssl:ro

depends_on:

- frontend
- backend

networks:

- app-network

frontend:

build: ./frontend

environment:

- REACT_APP_API_URL=http://localhost/api

networks:

- app-network

backend:

build: ./backend

environment:

- DATABASE_URL=postgres://user:pass@postgres:5432/mydb
- REDIS_URL=redis://redis:6379

depends_on:

- postgres
- redis

networks:

- app-network

postgres:

image: postgres:15-alpine

environment:

- POSTGRES_USER=user
- POSTGRES_PASSWORD=pass
- POSTGRES_DB=mydb

volumes:

- postgres_data:/var/lib/postgresql/data

networks:

- app-network

redis:

image: redis:7-alpine

volumes:

- redis_data:/data

networks:

- app-network

volumes:

postgres_data:

redis_data:

networks:

app-network:

driver: bridge

Development with Hot Reload



yaml

version: '3.8'

services:

app:

build:

context: .

dockerfile: Dockerfile.dev

volumes:

- ./src:/app/src # Mount source code
- /app/node_modules # Don't override node_modules

ports:

- "3000:3000"

environment:

- NODE_ENV=development

command: npm run dev

Multi-Stage Builds

Multi-stage builds let you use multiple FROM statements. Each stage can copy files from previous stages.

Why Multi-Stage Builds?

- Smaller final images (no build tools in production)
- Separate build and runtime dependencies
- Better security (fewer attack surfaces)
- Cleaner Dockerfiles

Basic Example



dockerfile

Stage 1: Build

FROM node:18 **AS** builder

WORKDIR /app

COPY package*.json ./

RUN npm install

COPY . .

RUN npm run build

Stage 2: Production

FROM node:18-alpine

WORKDIR /app

COPY --from=builder /app/dist ./dist

COPY --from=builder /app/node_modules ./node_modules

COPY package*.json ./

CMD ["node", "dist/server.js"]

Result: Final image only contains built files and runtime dependencies. Build tools stay in builder stage.

Real-World Examples

Go Application



dockerfile

Build stage

FROM golang:1.21 **AS** builder

WORKDIR /app

COPY go.* ./

RUN go mod download

COPY . .

RUN CGO_ENABLED=0 GOOS=linux go build -a -installsuffix cgo -o main .

Production stage

FROM alpine:latest

RUN apk --no-cache add ca-certificates

WORKDIR /root/

COPY --from=builder /app/main .

EXPOSE 8080

CMD ["/main"]

Size Comparison:

- With build tools: ~1GB
- Multi-stage: ~15MB

React Application



dockerfile

Build stage

FROM node:18 **AS** builder

WORKDIR /app

COPY package*.json ./

RUN npm ci

COPY . .

RUN npm run build

Production stage

FROM nginx:alpine

COPY --from=builder /app/build /usr/share/nginx/html

COPY nginx.conf /etc/nginx/conf.d/default.conf

EXPOSE 80

CMD ["nginx", "-g", "daemon off;"]

Python with Dependencies



dockerfile

```
# Build stage - compile dependencies
FROM python:3.11 AS builder
WORKDIR /app
COPY requirements.txt .
RUN pip install --user --no-cache-dir -r requirements.txt

# Production stage
FROM python:3.11-slim
WORKDIR /app
COPY --from=builder /root/.local /root/.local
COPY . .
ENV PATH=/root/.local/bin:$PATH
CMD ["python", "app.py"]
```

Java Application



dockerfile

```
# Build stage
FROM maven:3.9-eclipse-temurin-17 AS builder
WORKDIR /app
COPY pom.xml .
RUN mvn dependency:go-offline
COPY src ./src
RUN mvn package -DskipTests

# Production stage
FROM eclipse-temurin:17-jre-alpine
WORKDIR /app
COPY --from=builder /app/target/*.jar app.jar
EXPOSE 8080
ENTRYPOINT ["java", "-jar", "app.jar"]
```

Multiple Build Stages



dockerfile

```
# Dependencies stage
FROM node:18 AS dependencies
WORKDIR /app
COPY package*.json ./
RUN npm ci --only=production
RUN cp -R node_modules /tmp/node_modules
RUN npm ci

# Build stage
FROM node:18 AS builder
WORKDIR /app
COPY --from=dependencies /app/node_modules ./node_modules
COPY . .
RUN npm run build
RUN npm run test

# Production stage
FROM node:18-alpine
WORKDIR /app
COPY --from=dependencies /tmp/node_modules ./node_modules
COPY --from=builder /app/dist ./dist
COPY package*.json ./
CMD ["node", "dist/server.js"]
```

Named Stages and Build Targets



dockerfile

```
FROM node:18 AS base
WORKDIR /app
COPY package*.json ./
```

```
FROM base AS development
RUN npm install
COPY . .
CMD ["npm", "run", "dev"]
```

```
FROM base AS builder
RUN npm ci
COPY . .
RUN npm run build
```

```
FROM base AS production
RUN npm ci --only=production
COPY --from=builder /app/dist ./dist
CMD ["node", "dist/server.js"]
```

Build specific target:



bash

```
# Development
docker build --target development -t myapp:dev .
```

```
# Production
docker build --target production -t myapp:prod .
```

Docker Security

Security Best Practices

1. Don't Run as Root



dockerfile

Create user

FROM node:18-alpine

RUN addgroup -g 1001 -S nodejs && \

adduser -S nodejs -u 1001

USER nodejs



dockerfile

Python

FROM python:3.11-slim

RUN useradd -m -u 1000 appuser

USER appuser

2. Use Official Base Images



dockerfile

Good - official image

FROM node:18-alpine

Bad - unknown source

FROM someuser/node:latest

3. Scan Images for Vulnerabilities



bash

Docker Scout (built-in)

docker scout cves myimage:latest

docker scout recommendations myimage:latest

Trivy

trivy image myimage:latest

Snyk

snyk container **test** myimage:latest

4. Use Specific Tags



dockerfile

Bad

FROM node:latest

Good

FROM node:18.17.0-alpine3.18

5. Minimize Attack Surface



dockerfile

Use minimal base images

FROM alpine:3.18

FROM scratch *# Empty image (for compiled binaries)*

Remove unnecessary packages

RUN apt-get update && apt-get install -y \

curl \

&& rm -rf /var/lib/apt/lists/*

6. Don't Store Secrets in Images



dockerfile

```
# Bad - secret in image
ENV API_KEY=secret123

# Good - pass at runtime
docker run -e API_KEY=secret123 myimage

# Better - use Docker secrets or external secret manager
```

7. Use Read-Only Filesystem



bash

```
docker run --read-only \
  --tmpfs /tmp \
  --tmpfs /run \
  myimage
```

8. Limit Resources



bash

```
docker run \
  --memory="512m" \
  --cpus="1.0" \
  --pids-limit=100 \
  myimage
```

9. Use Security Options



bash

```
docker run \
  --security-opt=no-new-privileges:true \
  --cap-drop=ALL \
  --cap-add=NET_BIND_SERVICE \
  myimage
```

10. Enable Content Trust



bash

```
# Enable Docker Content Trust
export DOCKER_CONTENT_TRUST=1

# Now only signed images can be pulled
docker pull nginx:latest
```

Docker Bench Security

Automated security audit:



bash

```
docker run -it --net host --pid host --userns host --cap-add audit_control \
  -e DOCKER_CONTENT_TRUST=$DOCKER_CONTENT_TRUST \
  -v /var/lib:/var/lib \
  -v /var/run/docker.sock:/var/run/docker.sock \
  -v /usr/lib/systemd:/usr/lib/systemd \
  -v /etc:/etc --label docker_bench_security \
  docker/docker-bench-security
```

Secrets Management

Docker Secrets (Swarm Mode)



bash

Create secret

```
echo "my-secret-password" | docker secret create db_password -
```

Use in service

```
docker service create \
  --name db \
  --secret db_password \
  postgres:15
```

Access in container at /run/secrets/db_password

Environment Variables (Less Secure)



bash

Better than hardcoding

```
docker run -e API_KEY=$(cat api-key.txt) myimage
```

External Secret Managers

- HashiCorp Vault
- AWS Secrets Manager
- Azure Key Vault
- Google Secret Manager

Docker in Production

Production Checklist

Image Optimization



dockerfile

Use minimal base images

```
FROM alpine:3.18
```

Multi-stage builds

```
FROM node:18 AS builder
```

... build

```
FROM node:18-alpine
```

```
COPY --from=builder /app/dist ./dist
```

Combine RUN commands

```
RUN apk add --no-cache curl && \  
    rm -rf /var/cache/apk/*
```

Clean package manager cache

```
RUN apt-get update && apt-get install -y curl \  
    && rm -rf /var/lib/apt/lists/*
```

Health Checks



dockerfile

```
HEALTHCHECK --interval=30s --timeout=3s --retries=3 \  
    CMD curl -f http://localhost:8080/health || exit 1
```



yaml

docker-compose.yml

```
services:  
  api:  
    healthcheck:  
      test: ["CMD", "curl", "-f", "http://localhost:8080/health"]  
      interval: 30s  
      timeout: 10s  
      retries: 3  
      start_period: 40s
```

Logging



bash

JSON file driver (default)

```
docker run \
  --log-driver json-file \
  --log-opt max-size=10m \
  --log-opt max-file=3 \
  myimage
```

Syslog

```
docker run --log-driver syslog myimage
```

Fluentd

```
docker run --log-driver fluentd myimage
```

Disable logging (dangerous)

```
docker run --log-driver none myimage
```

Restart Policies



bash

```
docker run --restart=unless-stopped myimage
```

docker-compose.yml

```
services:
  app:
    restart: unless-stopped
```

Resource Limits



yaml

```
services:
  app:
    deploy:
      resources:
        limits:
          cpus: '2.0'
          memory: 1G
        reservations:
          cpus: '1.0'
          memory: 512M
```

Monitoring

View Resource Usage



bash

All containers

`docker stats`

Specific container

`docker stats container_name`

No streaming

`docker stats --no-stream`

Monitoring Solutions

- **Prometheus + Grafana:** Metrics and dashboards
- **ELK Stack:** Log aggregation
- **Datadog:** Full observability platform
- **New Relic:** Application monitoring

CI/CD Integration

GitHub Actions



yaml

name: Docker Build and Push

on:

push:

branches: [main]

jobs:

build:

runs-on: ubuntu-latest

steps:

- **uses:** actions/checkout@v3

- **name:** Set up Docker Buildx

uses: docker/setup-buildx-action@v2

- **name:** Login to DockerHub

uses: docker/login-action@v2

with:

username: \${ secrets.DOCKER_USERNAME }

password: \${ secrets.DOCKER_PASSWORD }

- **name:** Build and push

uses: docker/build-push-action@v4

with:

context: .

push: true

tags: myuser/myapp:latest

cache-from: type=registry,ref=myuser/myapp:buildcache

cache-to: type=registry,ref=myuser/myapp:buildcache,mode=max

GitLab CI



yaml

```
# .gitlab-ci.yml
```

```
build:
```

```
  image: docker:latest
```

```
  services:
```

```
    - docker:dind
```

```
  script:
```

```
    - docker login -u $CI_REGISTRY_USER -p $CI_REGISTRY_PASSWORD $CI_REGISTRY
```

```
    - docker build -t $CI_REGISTRY_IMAGE:$CI_COMMIT_SHA .
```

```
    - docker push $CI_REGISTRY_IMAGE:$CI_COMMIT_SHA
```

Docker Registry

Set Up Private Registry



```
bash
```

```
# Run registry
```

```
docker run -d \
```

```
  -p 5000:5000 \
```

```
  --name registry \
```

```
  -v registry_data:/var/lib/registry \
```

```
  registry:2
```

```
# Tag image
```

```
docker tag myapp:latest localhost:5000/myapp:latest
```

```
# Push
```

```
docker push localhost:5000/myapp:latest
```

```
# Pull
```

```
docker pull localhost:5000/myapp:latest
```

Registry with Authentication



```
bash
```

Create password file

```
htpasswd -Bc registry.password admin
```

Run with auth

```
docker run -d \  
-p 5000:5000 \  
--name registry \  
-v $(pwd)/registry.password:/auth/registry.password \  
-e REGISTRY_AUTH=htpasswd \  
-e REGISTRY_AUTH_HTPASSWD_PATH=/auth/registry.password \  
-e REGISTRY_AUTH_HTPASSWD_REALM="Registry Realm" \  
registry:2
```

Login

```
docker login localhost:5000
```

Advanced Topics

Docker Buildx

Enhanced build capabilities with BuildKit.



bash

Create builder

```
docker buildx create --name mybuilder --use
```

Build multi-platform images

```
docker buildx build \
  --platform linux/amd64,linux/arm64,linux/arm/v7 \
  -t myuser/myapp:latest \
  --push .
```

Build with cache

```
docker buildx build \
  --cache-from type=registry,ref=myuser/myapp:buildcache \
  --cache-to type=registry,ref=myuser/myapp:buildcache,mode=max \
  -t myuser/myapp:latest \
  --push .
```

BuildKit Features

Enable BuildKit:



bash

```
export DOCKER_BUILDKIT=1
```

Secret Mounts



dockerfile

```
# syntax=docker/dockerfile:1
```

```
FROM alpine
```

```
RUN --mount=type=secret,id=mysecret \
```

```
cat /run/secrets/mysecret
```



bash

```
docker build --secret id=mysecret,src=./secret.txt .
```

SSH Mounts



dockerfile

```
# syntax=docker/dockerfile:1
FROM alpine
RUN apk add git
RUN --mount=type=ssh \
    git clone git@github.com:user/private-repo.git
```



bash

```
docker build --ssh default .
```

Cache Mounts



dockerfile

```
# syntax=docker/dockerfile:1
FROM node:18
WORKDIR /app
COPY package*.json ./
RUN --mount=type=cache,target=/root/.npm \
    npm install
```

Docker Context

Manage multiple Docker environments.



bash

List contexts

`docker context ls`

Create context (SSH)

`docker context create remote \`
`--docker "host=ssh://user@remote-host"`

Create context (TCP)

`docker context create remote \`
`--docker "host=tcp://remote-host:2376"`

Use context

`docker context use remote`

Now all commands run on remote host

`docker ps`

Switch back

`docker context use default`

Docker System Commands



bash

Disk usage

`docker system df`

`docker system df -v` *# Verbose*

Clean everything

`docker system prune`

`docker system prune -a` *# Remove all unused images*

`docker system prune -a --volumes` *# Remove volumes too*

System info

`docker system info`

`docker info`

System events

`docker system events`

`docker events --since 1h`

Docker Save/Load vs Export/Import

Save/Load (Images)



bash

Save image with all layers and metadata

`docker save myapp:latest > myapp.tar`

`docker save myapp:latest | gzip > myapp.tar.gz`

Load image

`docker load < myapp.tar`

`docker load < myapp.tar.gz`

Export/Import (Containers)



bash

Export container filesystem (flattened, no history)

```
docker export container_name > container.tar
```

Import as image

```
docker import container.tar myapp:imported
```

```
cat container.tar | docker import - myapp:imported
```

Docker Plugins



bash

List plugins

```
docker plugin ls
```

Install plugin

```
docker plugin install vieux/sshfs
```

Enable/Disable

```
docker plugin disable vieux/sshfs
```

```
docker plugin enable vieux/sshfs
```

Remove plugin

```
docker plugin rm vieux/sshfs
```

Troubleshooting

Container Won't Start

Check Logs



bash

```
docker logs container_name
```

```
docker logs --tail 100 container_name
```

```
docker logs --since 10m container_name
```

Check Events



bash

```
docker events --since 10m
```

Inspect Container



bash

```
docker inspect container_name
docker inspect container_name | grep -A 10 State
```

Check Exit Code



bash

```
docker ps -a
# Look at STATUS column
# Exit (0) = normal exit
# Exit (1) = application error
# Exit (137) = killed by SIGKILL (often OOM)
# Exit (139) = segmentation fault
```

Networking Issues

Can't Connect to Container



bash

Check if container is running

`docker ps`

Check port mappings

`docker port container_name`

Check if process is listening inside container

`docker exec container_name netstat -tlnp`

Test from host

`curl localhost:8080`

Test from another container

`docker run --rm --network container:mycontainer alpine wget -O- localhost:8080`

Containers Can't Communicate



bash

Check if on same network

`docker network inspect network_name`

Check if DNS is working

`docker exec container1 nslookup container2`

`docker exec container1 ping container2`

Check firewall rules

`iptables -L -n`

Performance Issues

High CPU Usage



bash

Check stats

`docker stats`

Check processes inside container

`docker exec container_name top`

Limit CPU

`docker update --cpus 1.5 container_name`

High Memory Usage



bash

Check memory usage

`docker stats`

Check for memory leaks

`docker exec container_name ps aux`

Limit memory

`docker update --memory 512m container_name`

Slow Build Times



bash

Use BuildKit

`export DOCKER_BUILDKIT=1`

Use build cache

`docker build --cache-from myapp:latest .`

Order Dockerfile properly (stable layers first)

Use .dockerignore to exclude unnecessary files

Storage Issues

Out of Disk Space



bash

```
# Check disk usage
docker system df

# Remove unused data
docker system prune
docker system prune -a --volumes

# Remove specific items
docker image prune
docker container prune
docker volume prune
docker network prune
```

Volume Issues



bash

```
# List volumes
docker volume ls

# Inspect volume
docker volume inspect volume_name

# Find which container uses volume
docker ps -a --filter volume=volume_name

# Backup volume
docker run --rm \
-v volume_name:/data \
-v $(pwd):/backup \
alpine tar czf /backup/backup.tar.gz /data
```

Image Issues

Image Won't Build



bash

```
# Check Dockerfile syntax
docker build --no-cache .

# Build with verbose output
docker build --progress=plain .

# Check specific layer
docker build --target stage_name .
```

Can't Pull Image



bash

```
# Check network
ping registry.hub.docker.com

# Check auth
docker login

# Try different registry
docker pull registry.example.com/image:tag

# Check rate limits (Docker Hub)
TOKEN=$(curl "https://auth.docker.io/token?service=registry.docker.io&scope=repository:ratelimitpreview/test:pull" | jq .
curl --head -H "Authorization: Bearer $TOKEN" https://registry-1.docker.io/v2/ratelimitpreview/test/manifests/latest
```

Debug Running Container



bash

Get shell

`docker exec -it container_name bash`

`docker exec -it container_name sh` *# Alpine*

Check environment variables

`docker exec container_name env`

Check filesystem

`docker exec container_name ls -la /app`

Check running processes

`docker exec container_name ps aux`

Check network connections

`docker exec container_name netstat -tlnp`

Copy files out

`docker cp container_name:/app/logs/error.log ./`

View real-time logs

`docker logs -f container_name`

Common Error Messages

"Container already exists"



bash

`docker rm container_name`

or

`docker rm -f container_name` *# Force remove*

"Port already allocated"



bash

Find what's using the port

```
sudo lsof -i :8080
```

or

```
sudo netstat -tlnp | grep 8080
```

Use different port

```
docker run -p 8081:80 nginx
```

"No space left on device"



bash

```
docker system prune -a --volumes
```

"Cannot connect to Docker daemon"



bash

Start Docker

```
sudo systemctl start docker
```

Check if running

```
sudo systemctl status docker
```

Add user to docker group

```
sudo usermod -aG docker $USER
```

Log out and back in

"denied: access forbidden"



bash

```
# Login to registry
```

```
docker login
```

```
# Check credentials
```

```
cat ~/.docker/config.json
```

Interview Questions & Answers

Fundamentals

Q: What is Docker and why do we use it? A: Docker is a containerization platform that packages applications with their dependencies. We use it for consistency across environments, faster deployments, better resource utilization than VMs, easier scaling, and simplified dependency management. It solves the "works on my machine" problem.

Q: Difference between Docker image and container? A: An image is a read-only template (like a class in programming) containing the application and dependencies. A container is a running instance of an image (like an object). Images are built once, containers are created from images and can be started/stopped/deleted.

Q: What is a Dockerfile? A: A Dockerfile is a text file containing instructions to build a Docker image. It specifies the base image, copies files, installs dependencies, sets environment variables, and defines the command to run. It's like a recipe for creating an image.

Q: Difference between CMD and ENTRYPOINT? A: CMD provides default arguments that can be overridden when running the container. ENTRYPOINT defines the main command that always runs. Best practice is to use both: ENTRYPOINT for the main command, CMD for default arguments.



dockerfile

```
ENTRYPOINT ["python", "app.py"]
```

```
CMD ["--mode", "production"]
```

```
# docker run myimage -> python app.py --mode production
```

```
# docker run myimage --mode dev -> python app.py --mode dev
```

Q: Difference between COPY and ADD? A: COPY simply copies files from host to image. ADD does the same but also supports URLs and auto-extracts tar archives. Use COPY for simple file copying (preferred), use ADD only when you need its extra features.

Q: What is a Docker registry? A: A Docker registry is a storage system for Docker images. Docker Hub is the default public registry. Private registries include AWS ECR, Google GCR, Azure ACR, and self-hosted solutions like Harbor. You push images to registries and pull them when deploying.

Intermediate

Q: Explain Docker networking modes. A:

- Bridge (default): Isolated network, containers communicate via IP, need port publishing for host access

- Host: Container shares host network, no isolation, better performance
- None: No networking, completely isolated
- Custom bridge: User-defined network with DNS resolution, better isolation and control

Q: What are Docker volumes and why use them? A: Volumes persist data outside container lifecycle. Three types: named volumes (managed by Docker), bind mounts (host directory), tmpfs (memory). Use volumes for databases, user uploads, logs, or any data that should survive container restarts/deletions.

Q: What is Docker Compose? A: Docker Compose is a tool for defining multi-container applications in YAML. Instead of running multiple docker run commands, you define services, networks, and volumes in docker-compose.yml and manage everything with single commands like docker-compose up.

Q: Explain multi-stage builds. A: Multi-stage builds use multiple FROM statements in one Dockerfile. You build in one stage (with build tools) and copy artifacts to a minimal final stage. This keeps images small and secure by excluding build dependencies from production images.



dockerfile

```
FROM node:18 AS builder
WORKDIR /app
COPY . .
RUN npm install && npm run build
```

```
FROM node:18-alpine
COPY --from=builder /app/dist ./dist
CMD ["node", "dist/server.js"]
```

Q: How do containers communicate? A: Containers on the same custom network can reach each other by container name through Docker's embedded DNS server. Use docker network create to create custom networks. For external access, publish ports with -p flag.

Q: What is the difference between docker stop and docker kill? A: docker stop sends SIGTERM (graceful shutdown, allows cleanup) then SIGKILL after timeout. docker kill sends SIGKILL immediately (force termination). Always use stop in production to allow proper shutdown.

Advanced

Q: How do you optimize Docker images? A:

1. Use minimal base images (alpine, scratch)
2. Use multi-stage builds
3. Order Dockerfile layers properly (stable first, changing last)
4. Combine RUN commands to reduce layers
5. Use .dockerignore
6. Remove package manager caches
7. Use specific image tags, not latest

Q: Docker security best practices? A:

1. Don't run as root (USER instruction)

2. Use official base images
3. Scan images for vulnerabilities
4. Don't store secrets in images
5. Use read-only filesystems where possible
6. Limit container resources
7. Keep images updated
8. Use minimal base images
9. Enable Docker Content Trust

Q: How do you handle secrets in Docker? A: Never hardcode secrets in Dockerfiles or images. Options:

1. Docker secrets (Swarm mode)
2. Environment variables at runtime
3. External secret managers (Vault, AWS Secrets Manager)
4. Mount secrets as files at runtime
5. Use BuildKit secret mounts for build-time secrets

Q: Explain Docker layer caching. A: Each Dockerfile instruction creates a layer. Docker caches layers and reuses them if the instruction and its context haven't changed. If a layer changes, all subsequent layers are invalidated. Order instructions from least to most frequently changing for optimal caching.

Q: What is BuildKit? A: BuildKit is Docker's improved build backend with better performance, caching, and features. It supports parallel builds, build secrets, SSH mounts, cache mounts, and multi-platform builds. Enable with `DOCKER_BUILDKIT=1`.

Q: How would you debug a failing container? A:

1. Check logs: `docker logs container_name`
2. Check exit code: `docker ps -a`
3. Inspect container: `docker inspect container_name`
4. Get shell if possible: `docker exec -it container_name bash`
5. Check resource usage: `docker stats`
6. Review events: `docker events`
7. Test network connectivity
8. Check volume mounts

Q: Difference between docker-compose down and docker-compose stop? A: `stop` stops containers but keeps them. `down` stops AND removes containers, networks, and optionally volumes (with `-v`). Use `stop` for temporary pause, `down` for complete teardown.

Q: How do you limit container resources? A:



bash

```
docker run \  
  --memory="512m" \  
  --memory-swap="1g" \  
  --cpus="1.5" \  
  --pids-limit=100 \  
  myimage
```

Or in Compose:



yaml

```
services:
  app:
    deploy:
      resources:
        limits:
          cpus: '1.5'
          memory: 512M
```

Q: What happens when you run docker build? A:

- 1. Docker sends build context to daemon
- 2. Daemon reads Dockerfile
- 3. For each instruction, creates intermediate container
- 4. Runs instruction in container
- 5. Commits container to new layer
- 6. Removes intermediate container
- 7. Repeats for all instructions
- 8. Final layer becomes the image

Q: How do you implement health checks? A: In Dockerfile:



dockerfile

```
HEALTHCHECK --interval=30s --timeout=3s --retries=3 \
  CMD curl -f http://localhost/health || exit 1
```

Or in docker run:



bash

```
docker run --health-cmd="curl -f http://localhost/health" myimage
```

Q: What is the Docker storage driver? A: Storage driver manages the container's writable layer. overlay2 is the default and recommended. It uses union filesystem to layer images efficiently. Other drivers: aufs, btrfs, devicemapper, zfs.

Q: Explain Docker context. A: Docker context lets you manage multiple Docker environments (local, remote, different clusters) and switch between them. Useful for managing development, staging, and production environments.



bash

```
docker context create remote --docker "host=ssh://user@server"
docker context use remote
docker ps # Now running on remote host
```

Q: How would you deploy Docker in production? A:

- 1. Use orchestration (Kubernetes, Docker Swarm)
- 2. Implement health checks
- 3. Set restart policies
- 4. Configure logging
- 5. Limit resources
- 6. Use specific image tags
- 7. Implement monitoring
- 8. Set up CI/CD pipelines
- 9. Use secrets management
- 10. Regular security scans

Q: What is the difference between Union FS and Volume? A: Union FS (like overlay2) is how Docker layers images and container writable layers. It's copy-on-write and ephemeral. Volumes are for persistent data, stored outside the union filesystem, and survive container deletion.

Quick Reference Commands

Image Commands



bash

```
docker images           # List images
docker build -t name:tag . # Build image
docker pull image:tag    # Download image
docker push image:tag    # Upload image
docker rmi image:tag     # Remove image
docker image prune -a    # Remove unused images
docker history image:tag # Show image layers
docker inspect image:tag # Image details
docker save image > file.tar # Export image
docker load < file.tar    # Import image
```

Container Commands



bash

```
docker ps                # Running containers
docker ps -a             # All containers
docker run -d -p 8080:80 nginx # Run container
docker start container_name # Start container
docker stop container_name  # Stop container
docker restart container_name # Restart container
docker rm container_name    # Remove container
docker rm -f container_name # Force remove
docker exec -it container bash # Execute command
docker logs -f container_name # View logs
docker inspect container_name # Container details
docker stats              # Resource usage
docker cp file container:/path # Copy files
```

Network Commands



bash

```
docker network ls        # List networks
docker network create name # Create network
docker network connect net cont # Connect container
docker network disconnect net cont # Disconnect container
docker network inspect name # Network details
docker network rm name    # Remove network
docker network prune      # Remove unused networks
```

Volume Commands



bash


```
docker volume ls           # List volumes
docker volume create name  # Create volume
docker volume inspect name # Volume details
docker volume rm name      # Remove volume
docker volume prune        # Remove unused volumes
```

Docker Compose Commands



bash

```
docker-compose up           # Start services
docker-compose up -d        # Start detached
docker-compose down         # Stop and remove
docker-compose down -v      # Stop and remove volumes
docker-compose ps           # List services
docker-compose logs -f      # View logs
docker-compose exec service bash # Execute command
docker-compose build         # Build images
docker-compose restart      # Restart services
docker-compose stop         # Stop services
```

System Commands



bash

```
docker system df           # Disk usage
docker system prune        # Remove unused data
docker system prune -a     # Remove all unused data
docker system info         # System information
docker events              # System events
docker version             # Docker version
```

Best Practices Summary

Dockerfile

- Use official base images

- Use specific tags, not latest
- Order instructions from least to most frequently changing
- Combine RUN commands to reduce layers
- Use multi-stage builds
- Don't run as root
- Use .dockerignore
- Clean up in same RUN command

Security

- Don't store secrets in images
- Scan images regularly
- Use minimal base images
- Run as non-root user
- Keep images updated
- Limit container resources
- Use read-only filesystems where possible

Performance

- Use layer caching effectively
- Keep images small
- Use volumes for data
- Optimize for network I/O
- Use multi-stage builds
- Consider build cache strategies

Production

- Always use health checks
- Set restart policies
- Configure logging properly
- Monitor resource usage
- Use orchestration (Kubernetes/Swarm)
- Implement CI/CD
- Use specific image tags
- Regular backups of volumes

Common Patterns

Database Container



bash

```
docker run -d \
  --name postgres \
  -e POSTGRES_PASSWORD=secret \
  -e POSTGRES_DB=mydb \
  -v postgres_data:/var/lib/postgresql/data \
  -p 5432:5432 \
  --restart unless-stopped \
  postgres:15-alpine
```

Application with Database



yaml

```
version: '3.8'
services:
  app:
    build: .
    ports:
      - "3000:3000"
    environment:
      - DATABASE_URL=postgres://user:pass@db:5432/mydb
    depends_on:
      - db

  db:
    image: postgres:15-alpine
    environment:
      - POSTGRES_PASSWORD=pass
      - POSTGRES_USER=user
      - POSTGRES_DB=mydb
    volumes:
      - db_data:/var/lib/postgresql/data

volumes:
  db_data:
```

Development Environment



yaml

```
version: '3.8'

services:
  app:
    build:
      context: .
      dockerfile: Dockerfile.dev
    volumes:
      - ./src:/app/src
      - /app/node_modules
    ports:
      - "3000:3000"
    environment:
      - NODE_ENV=development
    command: npm run dev
```

Docker Swarm (Orchestration Basics)

Docker Swarm is Docker's native orchestration tool for managing clusters of Docker hosts.

Swarm Concepts

- **Node:** A Docker host in the swarm (manager or worker)
- **Manager:** Orchestrates and schedules services
- **Worker:** Executes containers
- **Service:** Definition of tasks to run on nodes
- **Task:** A single container running on a node
- **Stack:** Group of services defined in a Compose file

Basic Swarm Commands



bash

Initialize swarm

`docker swarm init`

Join swarm as worker

`docker swarm join --token TOKEN manager-ip:2377`

List nodes

`docker node ls`

Create service

`docker service create --name web --replicas 3 -p 80:80 nginx`

List services

`docker service ls`

Scale service

`docker service scale web=5`

Update service

`docker service update --image nginx:alpine web`

Remove service

`docker service rm web`

Leave swarm

`docker swarm leave`

Deploy Stack with Compose File



yaml

```
# docker-stack.yml
```

```
version: '3.8'
```

```
services:
```

```
  web:
```

```
    image: nginx:alpine
```

```
    ports:
```

```
      - "80:80"
```

```
    deploy:
```

```
      replicas: 3
```

```
      update_config:
```

```
        parallelism: 1
```

```
        delay: 10s
```

```
      restart_policy:
```

```
        condition: on-failure
```

```
      placement:
```

```
        constraints:
```

```
          - node.role == worker
```

```
  api:
```

```
    image: myapi:latest
```

```
    deploy:
```

```
      replicas: 2
```

```
      resources:
```

```
        limits:
```

```
          cpus: '0.5'
```

```
          memory: 512M
```

```
        reservations:
```

```
          cpus: '0.25'
```

```
          memory: 256M
```



```
bash
```

Deploy stack

```
docker stack deploy -c docker-stack.yml myapp
```

List stacks

```
docker stack ls
```

List services in stack

```
docker stack services myapp
```

Remove stack

```
docker stack rm myapp
```

Docker Performance Optimization

Build Performance

1. Use BuildKit



bash

```
export DOCKER_BUILDKIT=1
```

```
docker build .
```

2. Effective Layer Caching



dockerfile

Bad - cache invalidated on any code change

```
COPY . /app
```

```
RUN npm install
```

Good - dependencies cached separately

```
COPY package*.json /app/
```

```
RUN npm install
```

```
COPY . /app
```

3. Parallel Builds



dockerfile

```
# syntax=docker/dockerfile:1
FROM node:18 AS deps
COPY package*.json ./
RUN npm install

FROM node:18 AS build
COPY --from=deps /node_modules ./node_modules
COPY . .
RUN npm run build

# Both stages can build in parallel
```

4. Cache Mounts



dockerfile

```
# syntax=docker/dockerfile:1
FROM node:18
RUN --mount=type=cache,target=/root/.npm \
    npm install
```

Runtime Performance

1. Resource Limits



bash

```
docker run \
  --memory="512m" \
  --memory-reservation="256m" \
  --cpus="1.5" \
  --cpu-shares=1024 \
  myimage
```


2. Use Alpine Images



dockerfile

```
# Standard: ~900MB
FROM node:18

# Alpine: ~150MB
FROM node:18-alpine
```

3. Minimize Layers



dockerfile

```
# Bad - 3 layers
RUN apt-get update
RUN apt-get install curl
RUN apt-get install git

# Good - 1 layer
RUN apt-get update && \
    apt-get install -y curl git && \
    rm -rf /var/lib/apt/lists/*
```

4. Use .dockerignore



node_modules
.git
.gitignore
README.md
.env
.env.local
dist
build
coverage
.DS_Store
*.log

Network Performance

1. Use Host Network for Performance



bash

When you need maximum network performance

```
docker run --network host myimage
```

2. Optimize DNS



bash

```
docker run --dns 8.8.8.8 --dns 8.8.4.4 myimage
```

Storage Performance

1. Use Volumes (Not Bind Mounts)



bash

Better performance

```
docker run -v myvolume:/data myimage
```

Slower on Mac/Windows

```
docker run -v $(pwd):/data myimage
```

2. Use tmpfs for Temporary Data



bash

```
docker run --tmpfs /tmp:rw,size=100m myimage
```

Docker Anti-Patterns (What NOT to Do)

1. Running as Root



dockerfile

BAD

```
FROM ubuntu
COPY app /app
CMD ["/app"]
```

GOOD

```
FROM ubuntu
RUN useradd -m appuser
USER appuser
COPY app /app
CMD ["/app"]
```

2. Using latest Tag



dockerfile

```
# BAD
FROM node:latest
```

```
# GOOD
FROM node:18.17-alpine
```

3. Storing Data in Containers



bash

```
# BAD - data lost when container removed
docker run -d postgres

# GOOD
docker run -d -v postgres_data:/var/lib/postgresql/data postgres
```

4. Installing Unnecessary Packages



dockerfile

```
# BAD
RUN apt-get update && apt-get install -y \
    curl \
    wget \
    vim \
    nano \
    build-essential

# GOOD - only what you need
RUN apt-get update && apt-get install -y curl \
    && rm -rf /var/lib/apt/lists/*
```

5. Not Using .dockerignore

Without .dockerignore, you send node_modules, .git, etc. to build context.

6. Multiple Processes in One Container



dockerfile

```
# BAD
CMD service nginx start && service mysql start

# GOOD - use docker-compose for multiple services
```

7. Hardcoding Configuration



dockerfile

```
# BAD
ENV DATABASE_HOST=192.168.1.100
ENV API_KEY=secret123

# GOOD - pass at runtime
docker run -e DATABASE_HOST=db -e API_KEY=$API_KEY myimage
```

8. Not Using Health Checks



dockerfile

```
# BAD - no health check
FROM nginx

# GOOD
FROM nginx
HEALTHCHECK CMD curl -f http://localhost || exit 1
```

9. Ignoring Logs



bash

BAD - logs fill disk

```
docker run -d myapp
```

GOOD

```
docker run -d \  
  --log-opt max-size=10m \  
  --log-opt max-file=3 \  
  myapp
```

10. Not Cleaning Up



dockerfile

BAD

```
RUN apt-get update  
RUN apt-get install -y curl  
RUN wget https://example.com/file.tar.gz  
RUN tar xzf file.tar.gz
```

GOOD

```
RUN apt-get update && \  
  apt-get install -y curl && \  
  wget https://example.com/file.tar.gz && \  
  tar xzf file.tar.gz && \  
  rm file.tar.gz && \  
  apt-get clean && \  
  rm -rf /var/lib/apt/lists/*
```

Real-World Scenarios

Scenario 1: Microservices Architecture



yaml

version: '3.8'

services:

nginx:

image: nginx:alpine

ports:

- "80:80"
- "443:443"

volumes:

- ./nginx.conf:/etc/nginx/nginx.conf:ro
- ./ssl:/etc/nginx/ssl:ro

depends_on:

- frontend
- auth-service
- user-service
- product-service

networks:

- frontend

frontend:

build: ./frontend

environment:

- API_URL=http://api-gateway:3000

networks:

- frontend

api-gateway:

build: ./api-gateway

environment:

- AUTH_SERVICE=http://auth-service:3001
- USER_SERVICE=http://user-service:3002
- PRODUCT_SERVICE=http://product-service:3003

networks:

- frontend
- backend

auth-service:

build: ./services/auth

environment:

- DATABASE_URL=postgres://user:pass@postgres:5432/auth
- REDIS_URL=redis://redis:6379

depends_on:

- postgres
- redis

networks:

- backend

user-service:

build: ./services/user

environment:

- DATABASE_URL=postgres://user:pass@postgres:5432/users

depends_on:

- postgres

networks:

- backend

product-service:

build: ./services/product

environment:

- DATABASE_URL=postgres://user:pass@postgres:5432/products
- ELASTICSEARCH_URL=http://elasticsearch:9200

depends_on:

- postgres
- elasticsearch

networks:

- backend

postgres:

image: postgres:15-alpine

environment:

- POSTGRES_PASSWORD=secret

volumes:

- postgres_data:/var/lib/postgresql/data

networks:

- backend

redis:

image: redis:7-alpine

volumes:

- redis_data:/data

networks:

- backend

elasticsearch:

image: elasticsearch:8.10.0

environment:

- discovery.type=single-node
- ES_JAVA_OPTS=-Xms512m -Xmx512m

volumes:

- es_data:/usr/share/elasticsearch/data

networks:

- backend

volumes:

postgres_data:

redis_data:

es_data:

networks:

frontend:

backend:

Scenario 2: CI/CD Pipeline with Testing



yaml

```
# docker-compose.test.yml
```

```
version: '3.8'
```

```
services:
```

```
  app:
```

```
    build:
```

```
      context: .
```

```
      target: test
```

```
    environment:
```

```
      - NODE_ENV=test
```

```
      - DATABASE_URL=postgres://test:test@test-db:5432/testdb
```

```
    depends_on:
```

```
      - test-db
```

```
    command: npm test
```

```
  test-db:
```

```
    image: postgres:15-alpine
```

```
    environment:
```

```
      - POSTGRES_USER=test
```

```
      - POSTGRES_PASSWORD=test
```

```
      - POSTGRES_DB=testdb
```

```
    tmpfs:
```

```
      - /var/lib/postgresql/data
```



dockerfile

Multi-stage with test stage

```
FROM node:18-alpine AS base
WORKDIR /app
COPY package*.json ./
```

```
FROM base AS development
RUN npm install
COPY . .
CMD ["npm", "run", "dev"]
```

```
FROM base AS test
RUN npm install
COPY . .
CMD ["npm", "test"]
```

```
FROM base AS builder
RUN npm ci --only=production
COPY . .
RUN npm run build
```

```
FROM node:18-alpine AS production
WORKDIR /app
COPY --from=builder /app/dist ./dist
COPY --from=builder /app/node_modules ./node_modules
COPY package*.json ./
USER node
CMD ["node", "dist/server.js"]
```

Scenario 3: Development Environment with Hot Reload



yaml

version: '3.8'

services:

frontend:

build:

context: ./frontend

dockerfile: Dockerfile.dev

volumes:

- ./frontend/src:/app/src

- /app/node_modules

ports:

- "3000:3000"

environment:

- CHOKIDAR_USEPOLLING=true

- REACT_APP_API_URL=http://localhost:8000

command: npm start

backend:

build:

context: ./backend

dockerfile: Dockerfile.dev

volumes:

- ./backend:/app

- /app/node_modules

ports:

- "8000:8000"

environment:

- NODE_ENV=development

- DATABASE_URL=postgres://dev:dev@postgres:5432/devdb

depends_on:

- postgres

command: npm run dev

postgres:

image: postgres:15-alpine

environment:

- POSTGRES_USER=dev

- POSTGRES_PASSWORD=dev

- POSTGRES_DB=devdb

ports:

- "5432:5432"

volumes:

- postgres_dev:/var/lib/postgresql/data

volumes:

postgres_dev:

Scenario 4: Production with Monitoring



yaml

version: '3.8'

services:

app:

image: myapp:\${VERSION}

deploy:

replicas: 3

restart_policy:

condition: on-failure

max_attempts: 3

resources:

limits:

cpus: '1'

memory: 1G

reservations:

cpus: '0.5'

memory: 512M

environment:

- NODE_ENV=production

healthcheck:

test: ["CMD", "curl", "-f", "http://localhost:3000/health"]

interval: 30s

timeout: 10s

retries: 3

logging:

driver: "json-file"

options:

max-size: "10m"

max-file: "3"

networks:

- app-network

prometheus:

image: prom/prometheus

volumes:

- ./prometheus.yml:/etc/prometheus/prometheus.yml

- prometheus_data:/prometheus

ports:

- "9090:9090"

networks:

- app-network

grafana:
image: grafana/grafana
ports:
- "3001:3000"
environment:
- GF_SECURITY_ADMIN_PASSWORD=admin
volumes:
- grafana_data:/var/lib/grafana
networks:
- app-network

node-exporter:
image: prom/node-exporter
ports:
- "9100:9100"
networks:
- app-network

volumes:
prometheus_data:
grafana_data:

networks:
app-network:

Docker Tips & Tricks

1. Quick Container Inspection



bash

Get container IP

```
docker inspect -f '{{range .NetworkSettings.Networks}}{{.IPAddress}}{{end}}' container_name
```

Get container ports

```
docker inspect -f '{{range $p, $conf := .NetworkSettings.Ports}}{{$p}} -> {{(index $conf 0).HostPort}}{{end}}' container_name
```

Get environment variables

```
docker inspect -f '{{range .Config.Env}}{{println .}}{{end}}' container_name
```



2. Clean Up Everything



bash

Nuclear option - remove everything

```
docker system prune -a --volumes -f
```

Remove all stopped containers

```
docker container prune -f
```

Remove all dangling images

```
docker image prune -f
```

Remove all unused volumes

```
docker volume prune -f
```

3. Run Command in All Containers



bash

Restart all containers

```
docker ps -q | xargs docker restart
```

Stop all containers

```
docker ps -q | xargs docker stop
```

Remove all stopped containers

```
docker ps -aq | xargs docker rm
```

4. Copy Files Between Containers



bash

Container to container

```
docker cp container1:/path/file.txt /tmp/
```

```
docker cp /tmp/file.txt container2:/path/
```

One-liner

```
docker exec container1 cat /path/file.txt | docker exec -i container2 sh -c 'cat > /path/file.txt'
```

5. Monitor Multiple Containers



bash

Follow logs from multiple containers

```
docker-compose logs -f service1 service2
```

Stats for specific containers

```
docker stats container1 container2
```

Watch container status

```
watch 'docker ps --format "table {{.Names}}\t{{.Status}}\t{{.Ports}}"'
```

6. Quick Database Backup



bash

PostgreSQL

`docker exec postgres pg_dump -U user dbname > backup.sql`

MySQL

`docker exec mysql mysqldump -u user -ppassword dbname > backup.sql`

MongoDB

`docker exec mongo mongodump --out /tmp/backup`

`docker cp mongo:/tmp/backup ./backup`

7. Interactive Debugging



bash

Run ephemeral debug container

`docker run --rm -it \`
`--network container:target_container \`
`--pid container:target_container \`
`nicolaka/netshoot`

Attach debugger to running container

`docker run --rm -it \`
`--cap-add=SYS_PTRACE \`
`--pid=container:target \`
`alpine gdb -p 1`

8. Test DNS Resolution



bash

Check DNS from container

`docker run --rm alpine nslookup google.com`

Check DNS in specific network

`docker run --rm --network mynet alpine nslookup service_name`

9. Build Context Tricks



bash

```
# Build from URL
docker build github.com/user/repo
```

```
# Build from tar
docker build - < context.tar.gz
```

```
# Build with stdin Dockerfile
docker build -t myimage:latest - < Dockerfile
```

10. Container Forensics



bash

```
# Export container filesystem
docker export container_name > filesystem.tar
```

```
# Commit container to image (for debugging)
docker commit container_name debug-image
```

```
# Run container with different command
docker run --rm -it --entrypoint bash myimage
```

Docker Cheat Sheet

Essential Commands



bash

Images

<code>docker images</code>	<i># List</i>
<code>docker build -t name:tag .</code>	<i># Build</i>
<code>docker pull name:tag</code>	<i># Download</i>
<code>docker push name:tag</code>	<i># Upload</i>
<code>docker rmi name:tag</code>	<i># Remove</i>

Containers

<code>docker ps</code>	<i># List running</i>
<code>docker ps -a</code>	<i># List all</i>
<code>docker run -d -p 8080:80 nginx</code>	<i># Run</i>
<code>docker start/stop/restart name</code>	<i># Control</i>
<code>docker rm name</code>	<i># Remove</i>
<code>docker exec -it name bash</code>	<i># Execute</i>
<code>docker logs -f name</code>	<i># Logs</i>

Networks

<code>docker network ls</code>	<i># List</i>
<code>docker network create name</code>	<i># Create</i>
<code>docker network connect net cont</code>	<i># Connect</i>
<code>docker network inspect name</code>	<i># Inspect</i>

Volumes

<code>docker volume ls</code>	<i># List</i>
<code>docker volume create name</code>	<i># Create</i>
<code>docker volume inspect name</code>	<i># Inspect</i>
<code>docker volume rm name</code>	<i># Remove</i>

System

<code>docker system df</code>	<i># Disk usage</i>
<code>docker system prune</code>	<i># Clean up</i>
<code>docker info</code>	<i># System info</i>
<code>docker version</code>	<i># Version</i>

Compose

<code>docker-compose up -d</code>	<i># Start</i>
<code>docker-compose down</code>	<i># Stop</i>
<code>docker-compose logs -f</code>	<i># Logs</i>

```
docker-compose ps           # Status
docker-compose exec service bash  # Execute
```

Common Dockerfile Instructions



dockerfile

```
FROM image:tag           # Base image
WORKDIR /app             # Working directory
COPY src dest            # Copy files
RUN command              # Execute command
ENV KEY=value            # Environment variable
EXPOSE 8080              # Document port
VOLUME /data             # Mount point
USER username            # Switch user
CMD ["executable"]       # Default command
ENTRYPOINT ["executable"] # Main command
HEALTHCHECK CMD command  # Health check
```

Docker Run Options



bash

```
-d           # Detached mode
-it          # Interactive with TTY
-p 8080:80   # Port mapping
-e KEY=value # Environment variable
-v /host:/container # Volume mount
--name myapp # Container name
--network mynet # Custom network
--restart always # Restart policy
-m 512m       # Memory limit
--cpus 1.5    # CPU limit
--rm          # Remove after exit
-w /app       # Working directory
--user 1000:1000 # Run as user
```

Summary

You now have complete Docker mastery. Here's what you've learned:

Core Concepts

1. **Docker Basics:** Containers, images, Dockerfile, registry
2. **Difference from VMs:** Lighter, faster, shares kernel
3. **Layer System:** Caching, optimization, multi-stage builds

Practical Skills

4. **Image Building:** Efficient Dockerfiles, optimization techniques
5. **Container Management:** Run, stop, debug, monitor
6. **Networking:** Bridge, host, custom networks, DNS resolution
7. **Storage:** Volumes, bind mounts, data persistence
8. **Docker Compose:** Multi-container applications

Advanced Topics

9. **Security:** Run as non-root, scan images, secrets management
10. **Performance:** Caching, alpine images, resource limits
11. **Production:** Health checks, logging, monitoring, CI/CD
12. **Orchestration:** Docker Swarm basics
13. **Troubleshooting:** Debug techniques, common issues

Best Practices

14. Use specific tags, not latest
15. Order Dockerfile for optimal caching
16. Don't run as root
17. Use multi-stage builds
18. Keep images small
19. Never store secrets in images
20. Always use .dockerignore
21. Implement health checks
22. Clean up regularly

Interview Ready

You can now confidently answer questions about:

- Docker architecture and how it works
- Differences between images and containers
- Networking and storage strategies
- Security best practices
- Production deployment patterns
- Performance optimization
- Troubleshooting techniques

Next Steps

- Practice building real applications

- Learn Kubernetes (container orchestration at scale)
- Explore CI/CD pipelines with Docker
- Deep dive into security scanning
- Experiment with different deployment strategies

You're now a Docker expert! Go build amazing things!