

MongoDB Complete Guide: Zero to Hero

Table of Contents

- 1. Introduction to MongoDB
- 2. Installation & Setup
- 3. MongoDB Architecture
- 4. CRUD Operations
- 5. Data Modeling
- 6. Querying Deep Dive
- 7. Indexing
- 8. Aggregation Framework
- 9. Transactions
- 10. Replication
- 11. Sharding
- 12. Security
- 13. Performance Optimization
- 14. Backup & Restore
- 15. Interview Questions & Answers

1. Introduction to MongoDB

What is MongoDB?

MongoDB is a NoSQL document-oriented database. Instead of storing data in tables with rows and columns (like SQL), it stores data in flexible, JSON-like documents.

Key Features:

- Schema-less (flexible structure)
- Horizontally scalable
- High performance
- Rich query language
- Built-in replication and sharding

SQL vs NoSQL Comparison



SQL (MySQL)

MongoDB

Database

→ Database

Table

→ Collection

Row

→ Document

Column

→ Field

Table Join

→ Embedded documents or \$lookup

Primary Key

→ _id field (auto-generated)

When to Use MongoDB?

Good for:

- Rapidly changing schemas
- Hierarchical data storage
- Large scale data
- Real-time analytics
- Content management systems
- IoT applications
- Mobile apps

Not ideal for:

- Complex transactions (though now supported)
- Systems requiring complex joins
- Legacy systems built on SQL

2. Installation & Setup

Installation

Windows:



bash

```
# Download from MongoDB website
# Install MongoDB Community Server
# Add to PATH: C:\Program Files\MongoDB\Server\7.0\bin
```

Mac:



bash


```
brew tap mongodb/brew
brew install mongodb-community
brew services start mongodb-community
```

Linux (Ubuntu):



bash

```
wget -qO - https://www.mongodb.org/static/pgp/server-7.0.asc | sudo apt-key add -
echo "deb [ arch=amd64,arm64 ] https://repo.mongodb.org/apt/ubuntu focal/mongodb-org/7.0 multiverse" | sudo tee /etc/a
sudo apt-get update
sudo apt-get install -y mongodb-org
sudo systemctl start mongod
```



Starting MongoDB



bash

```
# Start MongoDB service
```

```
mongod
```

```
# Connect to MongoDB shell
```

```
mongosh
```

```
# Check MongoDB version
```

```
mongod --version
```

Basic Shell Commands



javascript

// Show all databases

```
show dbs
```

// Create or switch to database

```
use myDatabase
```

// Show current database

```
db
```

// Show collections

```
show collections
```

// Get database stats

```
db.stats()
```

// Drop database

```
db.dropDatabase()
```

3. MongoDB Architecture

Document Structure



javascript

```
{
  _id: ObjectId("507f1f77bcf86cd799439011"), // Unique identifier (auto-generated)
  name: "John Doe",
  age: 30,
  email: "john@example.com",
  address: {                                // Embedded document
    street: "123 Main St",
    city: "New York",
    zip: "10001"
  },
  hobbies: ["reading", "gaming", "coding"], // Array
  created_at: ISODate("2024-01-15T10:30:00Z")
}
```

Data Types in MongoDB



javascript

```
// String
{ name: "John" }

// Number (Integer, Long, Double)
{ age: 30, price: 99.99 }

// Boolean
{ isActive: true }

// Date
{ createdAt: new Date() }
{ createdAt: ISODate("2024-01-15") }

// Array
{ tags: ["mongodb", "database", "nosql"] }

// Embedded Document (Object)
{ address: { city: "NYC", zip: "10001" } }

// ObjectId (12-byte identifier)
{ _id: ObjectId("507f1f77bcf86cd799439011") }

// Null
{ middleName: null }

// Binary Data
{ file: BinData(0, "base64data") }

// Regular Expression
{ pattern: /^test/i }
```

Collections

Collections are groups of documents. They don't enforce a schema, so documents in the same collection can have different structures.



javascript

```
// Create collection explicitly
db.createCollection("users")

// Create with options
db.createCollection("logs", {
  capped: true,      // Fixed size
  size: 5242880,     // Max size in bytes (5MB)
  max: 5000          // Max documents
})
```

4. CRUD Operations

CREATE Operations

insertOne()



javascript

```
// Insert single document
db.users.insertOne({
  name: "Alice",
  age: 28,
  email: "alice@example.com",
  createdAt: new Date()
})

// Response
{
  acknowledged: true,
  insertedId: ObjectId("...")
}
```

insertMany()



javascript

// Insert multiple documents

```
db.users.insertMany([
  { name: "Bob", age: 35, email: "bob@example.com" },
  { name: "Charlie", age: 42, email: "charlie@example.com" },
  { name: "Diana", age: 29, email: "diana@example.com" }
])
```

// With ordered option (stops on first error if true)

```
db.users.insertMany(..., { ordered: false })
```

insert() (deprecated but still works)



javascript

```
db.users.insert({ name: "Eve", age: 31 })
```

READ Operations

find()



javascript

// Find all documents

```
db.users.find()
```

// Find with filter

```
db.users.find({ age: 28 })
```

// Find with multiple conditions (AND)

```
db.users.find({ age: 28, name: "Alice" })
```

// Find with OR

```
db.users.find({  
  $or: [  
    { age: 28 },  
    { name: "Bob" }  
  ]  
})
```

// Find with AND + OR

```
db.users.find({  
  age: { $gt: 25 },  
  $or: [  
    { name: "Alice" },  
    { name: "Bob" }  
  ]  
})
```

Comparison Operators



javascript

// \$eq (equal)

```
db.users.find({ age: { $eq: 30 } })
```

// \$ne (not equal)

```
db.users.find({ age: { $ne: 30 } })
```

// \$gt (greater than)

```
db.users.find({ age: { $gt: 30 } })
```

// \$gte (greater than or equal)

```
db.users.find({ age: { $gte: 30 } })
```

// \$lt (less than)

```
db.users.find({ age: { $lt: 30 } })
```

// \$lte (less than or equal)

```
db.users.find({ age: { $lte: 30 } })
```

// \$in (in array)

```
db.users.find({ age: { $in: [25, 30, 35] } })
```

// \$nin (not in array)

```
db.users.find({ age: { $nin: [25, 30, 35] } })
```

Logical Operators



javascript

```
// $and
db.users.find({
  $and: [
    { age: { $gt: 25 } },
    { age: { $lt: 40 } }
  ]
})
```

```
// $or
db.users.find({
  $or: [
    { name: "Alice" },
    { age: 35 }
  ]
})
```

```
// $not
db.users.find({ age: { $not: { $gt: 30 } } })
```

```
// $nor (not or)
db.users.find({
  $nor: [
    { age: { $lt: 25 } },
    { age: { $gt: 40 } }
  ]
})
```

Element Operators



javascript

```
// $exists (check if field exists)
db.users.find({ email: { $exists: true } })
```

```
// $type (check field type)
db.users.find({ age: { $type: "number" } })
db.users.find({ age: { $type: 16 } }) // 16 = 32-bit integer
```

Array Operators



javascript

```
// $all (array contains all elements)
db.users.find( { hobbies: { $all: ["reading", "coding"] } })

// $elemMatch (at least one element matches)
db.orders.find( {
  items: {
    $elemMatch: { price: { $gt: 50 }, quantity: { $gte: 2 } }
  }
})

// $size (array length)
db.users.find( { hobbies: { $size: 3 } })
```

Query Modifiers



javascript

```
// Projection (select specific fields)
db.users.find( {}, { name: 1, email: 1, _id: 0 } )

// Sort (1 = ascending, -1 = descending)
db.users.find().sort( { age: -1, name: 1 } )

// Limit
db.users.find().limit(5)

// Skip (pagination)
db.users.find().skip(10).limit(5)

// Count
db.users.find( { age: { $gt: 30 } }).count()
db.users.countDocuments( { age: { $gt: 30 } } )

// Distinct
db.users.distinct("age")
```

findOne()



javascript

```
// Returns first matching document
db.users.findOne({ name: "Alice" })

// With projection
db.users.findOne({ name: "Alice" }, { email: 1 })
```

Querying Embedded Documents



javascript

```
// Exact match
db.users.find({ address: { city: "NYC", zip: "10001" } })

// Dot notation (recommended)
db.users.find({ "address.city": "NYC" })
db.users.find({ "address.zip": "10001" })

// With operators
db.users.find({ "address.city": { $in: ["NYC", "LA"] } })
```

Querying Arrays



javascript

// Match array element

```
db.users.find({ hobbies: "reading" })
```

// Match entire array

```
db.users.find({ hobbies: ["reading", "coding"] })
```

// Array with condition

```
db.users.find({ hobbies: { $in: ["reading", "gaming"] } })
```

// Array element with dot notation

```
db.orders.find({ "items.0.name": "Laptop" }) // First element
```

UPDATE Operations

updateOne()



javascript

// Update first matching document

```
db.users.updateOne(
  { name: "Alice" },    // Filter
  { $set: { age: 29 } } // Update
)
```

// Response

```
{
  acknowledged: true,
  matchedCount: 1,
  modifiedCount: 1
}
```

updateMany()



javascript

// Update all matching documents

```
db.users.updateMany(  
  { age: { $lt: 30 } },  
  { $set: { status: "young" } }  
)
```

Update Operators



javascript

// \$set (set field value)

```
db.users.updateOne(  
  { name: "Alice" },  
  { $set: { age: 30, email: "newalice@example.com" } }  
)
```

// \$unset (remove field)

```
db.users.updateOne(  
  { name: "Alice" },  
  { $unset: { status: "" } }  
)
```

// \$inc (increment number)

```
db.users.updateOne(  
  { name: "Alice" },  
  { $inc: { age: 1 } }  
)
```

// \$mul (multiply number)

```
db.products.updateOne(  
  { name: "Laptop" },  
  { $mul: { price: 0.9 } } // 10% discount  
)
```

// \$min (update if new value is less)

```
db.products.updateOne(  
  { name: "Laptop" },  
  { $min: { price: 800 } }  
)
```

// \$max (update if new value is greater)

```
db.products.updateOne(  
  { name: "Laptop" },  
  { $max: { price: 1200 } }  
)
```

// \$rename (rename field)

```
db.users.updateMany(  
  {},  
  { $rename: { "email": "emailAddress" } }  
)
```

```
// $currentDate (set to current date)
db.users.updateOne(
  { name: "Alice" },
  { $currentDate: { lastModified: true } }
)
```

Array Update Operators



javascript

// \$push (add element to array)

```
db.users.updateOne(  
  { name: "Alice" },  
  { $push: { hobbies: "swimming" } }  
)
```

// \$push with \$each (add multiple elements)

```
db.users.updateOne(  
  { name: "Alice" },  
  { $push: { hobbies: { $each: ["cooking", "traveling"] } } }  
)
```

// \$push with \$sort

```
db.users.updateOne(  
  { name: "Alice" },  
  {  
    $push: {  
      scores: {  
        $each: [85, 92],  
        $sort: -1 // Sort descending  
      }  
    }  
  }  
)
```

// \$push with \$slice (limit array size)

```
db.users.updateOne(  
  { name: "Alice" },  
  {  
    $push: {  
      hobbies: {  
        $each: ["dancing"],  
        $slice: -5 // Keep last 5 elements  
      }  
    }  
  }  
)
```

// \$addToSet (add if not exists)

```
db.users.updateOne(  
  { name: "Alice" },
```

```
    { $addToSet: { hobbies: "reading" } }  
  )
```

// \$pop (remove first or last element)

```
db.users.updateOne(  
  { name: "Alice" },  
  { $pop: { hobbies: 1 } } // 1 = last, -1 = first  
)
```

// \$pull (remove matching elements)

```
db.users.updateOne(  
  { name: "Alice" },  
  { $pull: { hobbies: "gaming" } }  
)
```

// \$pull with condition

```
db.orders.updateOne(  
  { _id: ObjectId("...") },  
  { $pull: { items: { price: { $lt: 10 } } } }  
)
```

// \$pullAll (remove multiple values)

```
db.users.updateOne(  
  { name: "Alice" },  
  { $pullAll: { hobbies: ["gaming", "cooking"] } }  
)
```

// \$ (positional operator - update first match)

```
db.users.updateOne(  
  { name: "Alice", "hobbies": "reading" },  
  { $set: { "hobbies.$": "READING" } }  
)
```

// \$[] (update all array elements)

```
db.orders.updateOne(  
  { _id: ObjectId("...") },  
  { $set: { "items.$[].shipped": true } }  
)
```

// \$[element] (update filtered array elements)

```
db.orders.updateOne(  
  { $[element]: { $gt: 10 } },  
  { $set: { "$[element].shipped": true } }  
)
```

```
{ _id: ObjectId("...") },  
  { $set: { "items.$[elem].discount": 10 } },  
  { arrayFilters: [{ "elem.price": { $gt: 100 } }] }  
)
```

replaceOne()



javascript

```
// Replace entire document  
db.users.replaceOne(  
  { name: "Alice" },  
  {  
    name: "Alice Johnson",  
    age: 29,  
    email: "alice.j@example.com"  
  }  
)  
// Note: _id is preserved, all other fields are replaced
```

findOneAndUpdate()



javascript

```
// Update and return the document  
db.users.findOneAndUpdate(  
  { name: "Alice" },  
  { $inc: { age: 1 } },  
  { returnNewDocument: true } // Return updated document  
)
```

Upsert (Update or Insert)



javascript

// Insert if not exists, update if exists

```
db.users.updateOne(  
  { name: "Frank" },  
  { $set: { age: 40, email: "frank@example.com" } },  
  { upsert: true }  
)
```

DELETE Operations

deleteOne()



javascript

// Delete first matching document

```
db.users.deleteOne({ name: "Alice" })
```

// Response

```
{  
  acknowledged: true,  
  deletedCount: 1  
}
```

deleteMany()



javascript

// Delete all matching documents

```
db.users.deleteMany({ age: { $lt: 25 } })
```

// Delete all documents

```
db.users.deleteMany({})
```

findOneAndDelete()



javascript

// Delete and return the document

```
db.users.findOneAndDelete({ name: "Alice" })
```

Drop Collection



javascript

// Remove entire collection

```
db.users.drop()
```

5. Data Modeling

Embedded Documents (Denormalization)

When to use: One-to-One or One-to-Few relationships



javascript

// User with embedded address

```
{
  _id: ObjectId("..."),
  name: "John Doe",
  email: "john@example.com",
  address: {
    street: "123 Main St",
    city: "New York",
    state: "NY",
    zip: "10001"
  },
  phones: [
    { type: "home", number: "555-1234" },
    { type: "work", number: "555-5678" }
  ]
}
```

// Blog post with embedded comments

```
{
  _id: ObjectId("..."),
  title: "MongoDB Guide",
  content: "...",
  author: "John",
  comments: [
    { user: "Alice", text: "Great post!", date: ISODate("...") },
    { user: "Bob", text: "Very helpful", date: ISODate("...") }
  ]
}
```

Advantages:

- Better read performance (single query)
- Atomic updates
- Data locality

Disadvantages:

- Document size limit (16MB)
- Data duplication
- Harder to query embedded data

References (Normalization)

When to use: One-to-Many or Many-to-Many relationships



javascript

```
// Users collection
{
  _id: ObjectId("user1"),
  name: "John Doe",
  email: "john@example.com"
}

// Orders collection (referencing user)
{
  _id: ObjectId("order1"),
  user_id: ObjectId("user1"), // Reference
  total: 299.99,
  items: [...]
}

// Query with reference
db.orders.find({ user_id: ObjectId("user1") })
```

Many-to-Many Example:



javascript

```
// Students collection
{ _id: 1, name: "Alice", course_ids: [101, 102, 103] }
{ _id: 2, name: "Bob", course_ids: [101, 104] }

// Courses collection
{ _id: 101, name: "MongoDB Basics", student_ids: [1, 2] }
{ _id: 102, name: "Node.js", student_ids: [1] }
```

Hybrid Approach



javascript

// Store frequently accessed data embedded, reference for full details

```
{
  _id: ObjectId("order1"),
  user: {
    id: ObjectId("user1"),
    name: "John Doe",    // Embedded for quick access
    email: "john@example.com"
  },
  items: [
    {
      product_id: ObjectId("prod1"),
      name: "Laptop",    // Embedded
      price: 999,
      // Full product details in products collection
    }
  ]
}
```

Design Patterns

1. Attribute Pattern



javascript

// Instead of creating fields for each attribute

```
{  
  product: "Laptop",  
  color: "Silver",  
  size: "15-inch",  
  weight: "4.5 lbs"  
}
```

// Use attribute array for flexible attributes

```
{  
  product: "Laptop",  
  attributes: [  
    { k: "color", v: "Silver" },  
    { k: "size", v: "15-inch" },  
    { k: "weight", v: "4.5 lbs" }  
  ]  
}
```

// Create index on attributes.k and attributes.v

2. Bucket Pattern



javascript

// Instead of one document per measurement

```
{ sensor_id: 123, temp: 20.5, timestamp: ISODate("2024-01-01T10:00:00Z") }  
{ sensor_id: 123, temp: 20.7, timestamp: ISODate("2024-01-01T10:01:00Z") }
```

// Group measurements into buckets

```
{  
  sensor_id: 123,  
  date: ISODate("2024-01-01"),  
  measurements: [  
    { temp: 20.5, timestamp: ISODate("2024-01-01T10:00:00Z") },  
    { temp: 20.7, timestamp: ISODate("2024-01-01T10:01:00Z") },  
    // ... more measurements  
  ],  
  count: 60, // Number of measurements  
  avg_temp: 20.6  
}
```

3. Outlier Pattern



javascript

```
// Most products have few reviews, but some have thousands
// Normal product
{
  _id: ObjectId("prod1"),
  name: "Mouse",
  reviews: [
    { user: "Alice", rating: 5, text: "Great!" },
    { user: "Bob", rating: 4, text: "Good" }
  ]
}

// Popular product (use reference)
{
  _id: ObjectId("prod2"),
  name: "iPhone",
  review_count: 5000,
  has_overflow: true // Flag for outlier
}

// Separate reviews collection for outliers
{
  product_id: ObjectId("prod2"),
  user: "Charlie",
  rating: 5,
  text: "Amazing!"
}
```

4. Extended Reference Pattern



javascript

// Instead of just storing reference

```
{  
  order_id: ObjectId("..."),  
  customer_id: ObjectId("cust1")  
}
```

// Store frequently accessed fields

```
{  
  order_id: ObjectId("..."),  
  customer: {  
    id: ObjectId("cust1"),  
    name: "John Doe",      // Duplicate for performance  
    email: "john@example.com" // Duplicate for performance  
  }  
}
```

Schema Validation



javascript

// Create collection with validation rules

```
db.createCollection("users", {
  validator: {
    $jsonSchema: {
      bsonType: "object",
      required: ["name", "email", "age"],
      properties: {
        name: {
          bsonType: "string",
          description: "must be a string and is required"
        },
        email: {
          bsonType: "string",
          pattern: "^.+@.+$",
          description: "must be a valid email"
        },
        age: {
          bsonType: "int",
          minimum: 0,
          maximum: 150,
          description: "must be an integer between 0 and 150"
        },
        status: {
          enum: ["active", "inactive", "pending"],
          description: "can only be one of the enum values"
        }
      }
    }
  },
  validationLevel: "strict", // or "moderate"
  validationAction: "error" // or "warn"
})
```

// Add validation to existing collection

```
db.runCommand({
  collMod: "users",
  validator: { ... }
})
```

6. Querying Deep Dive

Regular Expressions



javascript

```
// Case-insensitive search
db.users.find({ name: { $regex: /john/i } })
db.users.find({ name: { $regex: "john", $options: "i" } })

// Starts with
db.users.find({ name: { $regex: /^john/i } })

// Ends with
db.users.find({ name: { $regex: /doe$/i } })

// Contains
db.users.find({ name: { $regex: /john/i } })
```

Text Search



javascript

// Create text index

```
db.articles.createIndex({ title: "text", content: "text" })
```

// Search for text

```
db.articles.find({ $text: { $search: "mongodb database" } })
```

// Search phrase

```
db.articles.find({ $text: { $search: "\"mongodb guide\"" } })
```

// Exclude words

```
db.articles.find({ $text: { $search: "mongodb -sql" } })
```

// Get text score

```
db.articles.find(  
  { $text: { $search: "mongodb" } },  
  { score: { $meta: "textScore" } }  
)<div data-bbox="23 404 342 424" data-label="Text">
```

Cursor Methods



javascript

// Get cursor

```
var cursor = db.users.find()
```

// Iterate cursor

```
cursor.forEach(doc => print(doc.name))
```

// Convert to array

```
var users = db.users.find().toArray()
```

// Check if cursor has next

```
cursor.hasNext()
```

// Get next document

```
cursor.next()
```

// Cursor count

```
cursor.count()
```

// Explain query execution

```
db.users.find({ age: { $gt: 30 } }).explain("executionStats")
```

Geospatial Queries



javascript

// Create 2dsphere index

```
db.places.createIndex({ location: "2dsphere" })
```

// Insert location data

```
db.places.insertOne({
  name: "Central Park",
  location: {
    type: "Point",
    coordinates: [-73.968285, 40.785091] // [longitude, latitude]
  }
})
```

// Find near a point

```
db.places.find({
  location: {
    $near: {
      $geometry: {
        type: "Point",
        coordinates: [-73.97, 40.78]
      },
      $maxDistance: 1000 // meters
    }
  }
})
```

// Find within area

```
db.places.find({
  location: {
    $geoWithin: {
      $centerSphere: [[-73.97, 40.78], 10 / 3963.2] // 10 miles radius
    }
  }
})
```

// Polygon search

```
db.places.find({
  location: {
    $geoWithin: {
      $geometry: {
        type: "Polygon",
        coordinates: [[
```

```
    [-73.99, 40.75],
    [-73.98, 40.75],
    [-73.98, 40.76],
    [-73.99, 40.76],
    [-73.99, 40.75]
  ]
}
}
}
})
```

7. Indexing

Why Indexes Matter

Without index: MongoDB scans every document (COLLSCAN) With index: MongoDB uses index tree structure (IXSCAN)



javascript

```
// Without index - slow on large collections
db.users.find({ email: "alice@example.com" }) // Scans all documents

// Create index
db.users.createIndex({ email: 1 })

// With index - fast
db.users.find({ email: "alice@example.com" }) // Uses index
```

Single Field Index



javascript

// Ascending index

```
db.users.createIndex({ age: 1 })
```

// Descending index

```
db.users.createIndex({ age: -1 })
```

// Check existing indexes

```
db.users.getIndexes()
```

// Drop index

```
db.users.dropIndex({ age: 1 })
```

```
db.users.dropIndex("age_1") // By name
```

// Drop all indexes (except _id)

```
db.users.dropIndexes()
```

Compound Index



javascript

// Index on multiple fields

```
db.users.createIndex({ age: 1, name: 1 })
```

// Order matters!

// Good for: { age: 30, name: "John" }

// Good for: { age: 30 }

// NOT optimized for: { name: "John" }

// Best practice: Equality, Sort, Range (ESR)

```
db.orders.createIndex({  
  status: 1,    // Equality  
  created_at: -1, // Sort  
  total: 1      // Range  
})
```

Index Types

1. Unique Index



javascript

```
// Prevent duplicate values
db.users.createIndex({ email: 1 }, { unique: true })

// Compound unique
db.users.createIndex({ email: 1, phone: 1 }, { unique: true })
```

2. Sparse Index



javascript

```
// Index only documents that have the field
db.users.createIndex({ phone: 1 }, { sparse: true })

// Documents without 'phone' field won't be in index
```

3. TTL Index (Time To Live)



javascript

```
// Auto-delete documents after specified time
db.sessions.createIndex(
  { createdAt: 1 },
  { expireAfterSeconds: 3600 } // Delete after 1 hour
)

// Insert with createdAt
db.sessions.insertOne({
  user_id: "user123",
  createdAt: new Date()
})
```

4. Partial Index



javascript

// Index only documents matching filter

```
db.orders.createIndex(  
  { customer_id: 1, total: 1 },  
  { partialFilterExpression: { status: "active" } }  
)
```

// Only active orders are indexed

5. Text Index



javascript

// For text search

```
db.articles.createIndex({ content: "text" })
```

// Multiple fields

```
db.articles.createIndex({  
  title: "text",  
  content: "text"  
})
```

// With weights (importance)

```
db.articles.createIndex(  
  {  
    title: "text",  
    content: "text"  
  },  
  {  
    weights: {  
      title: 10,  
      content: 5  
    }  
  }  
)
```

6. Wildcard Index



javascript

// Index all fields

```
db.products.createIndex({ "$**": 1 })
```

// Index all subfields of specific field

```
db.products.createIndex({ "attributes.$**": 1 })
```

7. Hashed Index



javascript

// For sharding

```
db.users.createIndex({ _id: "hashed" })
```

8. 2dsphere Index (Geospatial)



javascript

```
db.places.createIndex({ location: "2dsphere" })
```

Index Properties



javascript

// Background building (don't block operations)

```
db.users.createIndex({ age: 1 }, { background: true })
```

// Unique

```
db.users.createIndex({ email: 1 }, { unique: true })
```

// Sparse

```
db.users.createIndex({ phone: 1 }, { sparse: true })
```

// Name

```
db.users.createIndex({ age: 1 }, { name: "age_ascending" })
```

// Partial

```
db.orders.createIndex(  
  { total: 1 },  
  { partialFilterExpression: { total: { $gt: 100 } } }  
)
```

// TTL

```
db.logs.createIndex({ timestamp: 1 }, { expireAfterSeconds: 86400 })
```

// Case insensitive (collation)

```
db.users.createIndex(  
  { name: 1 },  
  { collation: { locale: "en", strength: 2 } }  
)
```

Index Analysis



javascript

// Explain query plan

```
db.users.find({ age: 30 }).explain("executionStats")
```

// Key fields in explain output:

// - executionTimeMillis: Time taken

// - totalDocsExamined: Documents scanned

// - totalKeysExamined: Index entries scanned

// - stage: IXSCAN (index) or COLLSCAN (collection)

// Index statistics

```
db.users.aggregate([ { $indexStats: {} } ])
```

// Get index sizes

```
db.users.stats().indexSizes
```

Covered Queries



javascript

// Query covered entirely by index (no document lookup)

```
db.users.createIndex({ name: 1, age: 1 })
```

// This query is covered (returns only indexed fields)

```
db.users.find(  
  { name: "John" },  
  { name: 1, age: 1, _id: 0 }  
)
```

// executionStats shows totalDocsExamined: 0

Index Best Practices

1. Create indexes for frequently queried fields
 2. Use compound indexes wisely (ESR rule)
 3. Limit number of indexes (write performance impact)
 4. Use covered queries when possible
 5. Monitor index usage with \$indexStats
 6. Delete unused indexes
 7. Consider partial indexes for large collections
-

8. Aggregation Framework

The aggregation pipeline processes documents through stages, each transforming the data.

Basic Pipeline



javascript

```
db.collection.aggregate([
  { $match: { ... } },    // Stage 1: Filter
  { $group: { ... } },    // Stage 2: Group
  { $sort: { ... } },     // Stage 3: Sort
  { $project: { ... } }   // Stage 4: Shape output
])
```

Common Aggregation Stages

\$match (Filter)



javascript

```
// Filter documents (like find)
db.orders.aggregate([
  { $match: { status: "completed" } }
])

// With operators
db.orders.aggregate([
  { $match: {
    total: { $gt: 100 },
    status: "completed"
  }}
])
```

\$project (Shape Output)



javascript

// Select/exclude fields

```
db.users.aggregate([  
  { $project: {  
    name: 1,  
    email: 1,  
    _id: 0  
  }}  
])
```

// Compute new fields

```
db.orders.aggregate([  
  { $project: {  
    total: 1,  
    tax: { $multiply: ["$total", 0.1] },  
    totalWithTax: { $add: ["$total", { $multiply: ["$total", 0.1] }] }  
  }}  
])
```

// String operations

```
db.users.aggregate([  
  { $project: {  
    fullName: { $concat: ["$firstName", " ", "$lastName"] },  
    upperName: { $toUpper: "$name" },  
    emailDomain: { $arrayElemAt: [{ $split: ["$email", "@"] }, 1] }  
  }}  
])
```

\$group (Aggregation)



javascript

// Count documents

```
db.orders.aggregate([
  { $group: {
    _id: null,
    count: { $sum: 1 }
  }}
])
```

// Group by field

```
db.orders.aggregate([
  { $group: {
    _id: "$status",
    count: { $sum: 1 }
  }}
])
```

// Multiple grouping fields

```
db.orders.aggregate([
  { $group: {
    _id: {
      status: "$status",
      year: { $year: "$created_at" }
    },
    count: { $sum: 1 }
  }}
])
```

// Aggregation operators

```
db.orders.aggregate([
  { $group: {
    _id: "$customer_id",
    totalSpent: { $sum: "$total" },    // Sum
    avgOrder: { $avg: "$total" },     // Average
    maxOrder: { $max: "$total" },     // Maximum
    minOrder: { $min: "$total" },     // Minimum
    firstOrder: { $first: "$created_at" }, // First value
    lastOrder: { $last: "$created_at" }, // Last value
    orders: { $push: "$_id" },        // Array of values
    uniqueStatuses: { $addToSet: "$status" } // Unique values
  }}
])
```

```
  }}  
  )
```

\$sort



javascript

```
// Sort results  
db.orders.aggregate([  
  { $group: {  
    _id: "$customer_id",  
    totalSpent: { $sum: "$total" }  
  }},  
  { $sort: { totalSpent: -1 } } // Descending  
])
```

\$limit and \$skip



javascript

```
// Pagination  
db.orders.aggregate([  
  { $sort: { created_at: -1 } },  
  { $skip: 20 },  
  { $limit: 10 }  
])
```

\$unwind (Deconstruct Array)



javascript

// Deconstruct array field

```
db.orders.aggregate([  
  { $unwind: "$items" }  
])
```

// Before: { _id: 1, items: ["A", "B", "C"] }

// After: { _id: 1, items: "A" }

// { _id: 1, items: "B" }

// { _id: 1, items: "C" }

// Preserve empty arrays

```
db.orders.aggregate([  
  { $unwind: {  
    path: "$items",  
    preserveNullAndEmptyArrays: true  
  }}  
])
```

\$lookup (Join)



javascript

// Left outer join

```
db.orders.aggregate([
  {
    $lookup: {
      from: "customers",      // Collection to join
      localField: "customer_id", // Field from orders
      foreignField: "_id",     // Field from customers
      as: "customerInfo"      // Output array field
    }
  }
])
```

// Unwind result

```
db.orders.aggregate([
  {
    $lookup: {
      from: "customers",
      localField: "customer_id",
      foreignField: "_id",
      as: "customer"
    }
  },
  { $unwind: "$customer" }
])
```

// Complex lookup with pipeline

```
db.orders.aggregate([
  {
    $lookup: {
      from: "products",
      let: { order_items: "$items" },
      pipeline: [
        { $match: {
            $expr: { $in: ["$_id", "$order_items"] }
          }
        },
        { $project: { name: 1, price: 1 } }
      ],
      as: "productDetails"
    }
  }
])
```

```
}  
])
```

\$addFields



javascript

```
// Add new fields without removing existing ones  
db.orders.aggregate([  
  { $addFields: {  
    totalWithTax: { $multiply: ["$total", 1.1] },  
    year: { $year: "$created_at" }  
  }}  
])
```

\$replaceRoot



javascript

```
// Replace document with embedded document  
db.orders.aggregate([  
  { $replaceRoot: { newRoot: "$customer" } }  
])
```

\$out



javascript

// Write results to new collection

```
db.orders.aggregate([
  { $match: { status: "completed" } },
  { $group: {
    _id: "$customer_id",
    total: { $sum: "$total" }
  }},
  { $out: "customer_totals" } // Creates new collection
])
```

\$merge



javascript

// Merge results into existing collection

```
db.orders.aggregate([
  { $group: {
    _id: "$customer_id",
    totalSpent: { $sum: "$total" }
  }},
  { $merge: {
    into: "customers",
    on: "_id",
    whenMatched: "merge", // or "replace", "keepExisting", "fail"
    whenNotMatched: "discard" // or "insert"
  }}
])
```

\$bucket



javascript

// Group by ranges

```
db.users.aggregate([
  {
    $bucket: {
      groupBy: "$age",
      boundaries: [0, 18, 30, 50, 100],
      default: "Other",
      output: {
        count: { $sum: 1 },
        users: { $push: "$name" }
      }
    }
  }
])
```

\$facet



javascript

// Multiple parallel pipelines

```
db.products.aggregate([
  {
    $facet: {
      categoryCounts: [
        { $group: { _id: "$category", count: { $sum: 1 } } }
      ],
      priceStats: [
        { $group: {
          _id: null,
          avgPrice: { $avg: "$price" },
          maxPrice: { $max: "$price" }
        }}
      ],
      topProducts: [
        { $sort: { sales: -1 } },
        { $limit: 5 }
      ]
    }
  }
])
```

Real-World Examples

Example 1: Sales Report



javascript

```
db.orders.aggregate([
  // Filter completed orders
  { $match: {
    status: "completed",
    created_at: { $gte: ISODate("2024-01-01") }
  }},

  // Unwind items array
  { $unwind: "$items" },

  // Group by product
  { $group: {
    _id: "$items.product_id",
    totalQuantity: { $sum: "$items.quantity" },
    totalRevenue: { $sum: { $multiply: ["$items.quantity", "$items.price"] } },
    orderCount: { $sum: 1 }
  }},

  // Lookup product details
  { $lookup: {
    from: "products",
    localField: "_id",
    foreignField: "_id",
    as: "product"
  }},
  { $unwind: "$product" },

  // Shape output
  { $project: {
    _id: 0,
    productName: "$product.name",
    totalQuantity: 1,
    totalRevenue: 1,
    orderCount: 1,
    avgOrderValue: { $divide: ["$totalRevenue", "$orderCount"] }
  }},

  // Sort by revenue
  { $sort: { totalRevenue: -1 } },

  // Top 10
```

```
{ $limit: 10 }
```

```
)
```

Example 2: Customer Analysis



javascript

```

db.orders.aggregate([
  // Group by customer
  { $group: {
    _id: "$customer_id",
    totalOrders: { $sum: 1 },
    totalSpent: { $sum: "$total" },
    avgOrderValue: { $avg: "$total" },
    firstOrder: { $min: "$created_at" },
    lastOrder: { $max: "$created_at" }
  }},

  // Calculate days since first order
  { $addFields: {
    daysSinceFirst: {
      $divide: [
        { $subtract: [new Date(), "$firstOrder"] },
        1000 * 60 * 60 * 24
      ]
    }
  }},

  // Customer lifetime value per day
  { $addFields: {
    clvPerDay: { $divide: ["$totalSpent", "$daysSinceFirst"] }
  }},

  // Join with customer details
  { $lookup: {
    from: "customers",
    localField: "_id",
    foreignField: "_id",
    as: "customer"
  }},
  { $unwind: "$customer" },

  // Filter high-value customers
  { $match: { totalSpent: { $gt: 1000 } } },

  // Sort by CLV

```

```
{ $sort: { clvPerDay: -1 } }  
])
```

Example 3: Time Series Analysis



javascript

```
db.orders.aggregate([
  // Group by year and month
  { $group: {
    _id: {
      year: { $year: "$created_at" },
      month: { $month: "$created_at" }
    },
    revenue: { $sum: "$total" },
    orders: { $sum: 1 },
    avgOrderValue: { $avg: "$total" }
  }},

  // Sort chronologically
  { $sort: { "_id.year": 1, "_id.month": 1 } },

  // Calculate month-over-month growth
  { $setWindowFields: {
    sortBy: { "_id.year": 1, "_id.month": 1 },
    output: {
      previousRevenue: {
        $shift: {
          output: "$revenue",
          by: -1
        }
      }
    }
  }},

  { $addFields: {
    growthPercent: {
      $multiply: [
        { $divide: [
          { $subtract: ["$revenue", "$previousRevenue"] },
          "$previousRevenue"
        ]},
        100
      ]
    }
  }},

  // Format output
```

```
{ $project: {
  _id: 0,
  period: { $concat: [
    { $toString: "$_id.year" },
    "-",
    { $toString: "$_id.month" }
  ] },
  revenue: { $round: ["$revenue", 2] },
  orders: 1,
  avgOrderValue: { $round: ["$avgOrderValue", 2] },
  growthPercent: { $round: ["$growthPercent", 2] }
}}
```

9. Transactions

MongoDB supports ACID transactions for multi-document operations.

Single Document Transactions



javascript

// Single document operations are ALWAYS atomic

```
db.accounts.updateOne(
  { _id: "account1" },
  { $inc: { balance: -100 } }
)
```

// This is atomic - either completes fully or not at all

Multi-Documnt Transactions



javascript

// Start a session

```
const session = db.getMongo().startSession()
```

// Start transaction

```
session.startTransaction()
```

```
try {
```

```
    const accounts = session.getDatabase("bank").accounts
```

// Debit from account1

```
accounts.updateOne(
    { _id: "account1" },
    { $inc: { balance: -100 } },
    { session }
)
```

// Credit to account2

```
accounts.updateOne(
    { _id: "account2" },
    { $inc: { balance: 100 } },
    { session }
)
```

// Commit transaction

```
session.commitTransaction()
print("Transaction committed")
```

```
} catch (error) {
```

// Abort transaction on error

```
session.abortTransaction()
print("Transaction aborted: " + error)
```

```
} finally {
```

```
    session.endSession()
}
```

Transaction with Node.js Driver



javascript

```

const { MongoClient } = require('mongodb')

async function transferMoney() {
  const client = new MongoClient('mongodb://localhost:27017')
  await client.connect()

  const session = client.startSession()

  try {
    await session.withTransaction(async () => {
      const accounts = client.db('bank').collection('accounts')

      await accounts.updateOne(
        { _id: 'account1' },
        { $inc: { balance: -100 } },
        { session }
      )

      await accounts.updateOne(
        { _id: 'account2' },
        { $inc: { balance: 100 } },
        { session }
      )
    })

    console.log('Transaction successful')
  } catch (error) {
    console.log('Transaction failed:', error)
  } finally {
    await session.endSession()
    await client.close()
  }
}

```

Transaction Best Practices

1. **Keep transactions short** (ideally < 1 second)
 2. **Limit operations per transaction** (< 1000 documents)
 3. **Use appropriate read/write concerns**
 4. **Handle retries** (transient errors)
 5. **Consider using single-document operations when possible**
-

10. Replication

Replication provides redundancy and high availability.

Replica Set Architecture



Primary Node (Read/Write)

|

-- Secondary Node (Read only)

-- Secondary Node (Read only)

-- Arbiter (Voting only, no data)

Setting Up Replica Set



bash

Start three mongod instances

mongod --replSet rs0 --port 27017 --dbpath /data/db1

mongod --replSet rs0 --port 27018 --dbpath /data/db2

mongod --replSet rs0 --port 27019 --dbpath /data/db3



javascript

// Connect to first instance

```
mongosh --port 27017
```

// Initiate replica set

```
rs.initiate({  
  _id: "rs0",  
  members: [  
    { _id: 0, host: "localhost:27017" },  
    { _id: 1, host: "localhost:27018" },  
    { _id: 2, host: "localhost:27019" }  
  ]  
})
```

// Check replica set status

```
rs.status()
```

// Check replica set configuration

```
rs.conf()
```

// Add member

```
rs.add("localhost:27020")
```

// Remove member

```
rs.remove("localhost:27020")
```

Replica Set Member Types



javascript

```

// Priority (likelihood of becoming primary)
rs.reconfig({
  _id: "rs0",
  members: [
    { _id: 0, host: "localhost:27017", priority: 2 }, // Preferred primary
    { _id: 1, host: "localhost:27018", priority: 1 },
    { _id: 2, host: "localhost:27019", priority: 0 } // Never primary
  ]
})

// Hidden member (doesn't receive client reads)
{ _id: 3, host: "localhost:27020", hidden: true, priority: 0 }

// Delayed member (for backup/disaster recovery)
{ _id: 4, host: "localhost:27021", slaveDelay: 3600, priority: 0 } // 1 hour delay

// Arbiter (voting only, no data)
rs.addArb("localhost:27022")

```

Read Preferences



javascript

```

// Read from primary (default)
db.users.find().readPref("primary")

// Read from primary preferred (primary if available, else secondary)
db.users.find().readPref("primaryPreferred")

// Read from secondary
db.users.find().readPref("secondary")

// Read from secondary preferred
db.users.find().readPref("secondaryPreferred")

// Read from nearest (lowest latency)
db.users.find().readPref("nearest")

```

Write Concerns



javascript

// Write acknowledged by primary only (default)

```
db.users.insertOne(  
  { name: "John" },  
  { writeConcern: { w: 1 } }  
)
```

// Write acknowledged by majority

```
db.users.insertOne(  
  { name: "John" },  
  { writeConcern: { w: "majority" } }  
)
```

// Write acknowledged by specific number

```
db.users.insertOne(  
  { name: "John" },  
  { writeConcern: { w: 2 } }  
)
```

// Write with timeout

```
db.users.insertOne(  
  { name: "John" },  
  { writeConcern: { w: "majority", wtimeout: 5000 } }  
)
```

// Journalled write

```
db.users.insertOne(  
  { name: "John" },  
  { writeConcern: { w: 1, j: true } }  
)
```

Read Concerns



javascript

```
// Local (default - reads latest data)
db.users.find().readConcern("local")
```

```
// Available (no guarantee of durability)
db.users.find().readConcern("available")
```

```
// Majority (reads data acknowledged by majority)
db.users.find().readConcern("majority")
```

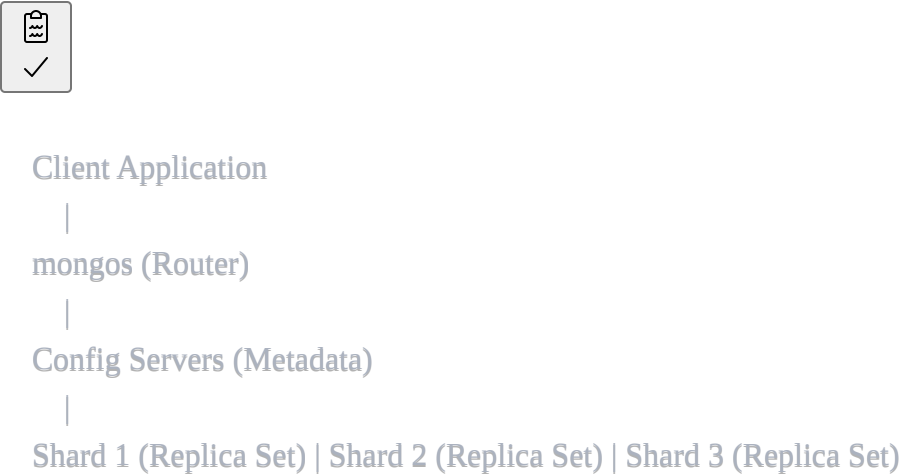
```
// Linearizable (reads reflect all prior majority writes)
db.users.find().readConcern("linearizable")
```

```
// Snapshot (for transactions)
session.startTransaction({ readConcern: { level: "snapshot" } })
```

11. Sharding

Sharding distributes data across multiple machines for horizontal scaling.

Sharding Architecture



Setting Up Sharding

1. Start Config Servers

```
bash
```

```
mongod --configsvr --replSet configRS --port 27019 --dbpath /data/configdb
```

2. Start Shards



bash

```
mongod --shardsvr --replSet shard1RS --port 27018 --dbpath /data/shard1
```

```
mongod --shardsvr --replSet shard2RS --port 27020 --dbpath /data/shard2
```

3. Start mongos



bash

```
mongos --configdb configRS/localhost:27019 --port 27017
```

4. Add Shards



javascript

```
// Connect to mongos
```

```
mongosh --port 27017
```

```
// Add shards
```

```
sh.addShard("shard1RS/localhost:27018")
```

```
sh.addShard("shard2RS/localhost:27020")
```

```
// Check status
```

```
sh.status()
```

5. Enable Sharding



javascript

// Enable sharding on database

```
sh.enableSharding("myDatabase")
```

// Shard collection

```
sh.shardCollection("myDatabase.users", { user_id: 1 })
```

// Check sharding status

```
db.users.getShardDistribution()
```

Shard Key Selection



javascript

// Hashed shard key (even distribution)

```
sh.shardCollection("myDatabase.users", { _id: "hashed" })
```

// Range-based shard key

```
sh.shardCollection("myDatabase.orders", { customer_id: 1, created_at: 1 })
```

// Compound shard key

```
sh.shardCollection("myDatabase.products", { category: 1, product_id: 1 })
```

Shard Key Best Practices

Good Shard Keys:

- High cardinality (many unique values)
- Even distribution
- Query isolation (most queries include shard key)

Bad Shard Keys:

- Monotonically increasing (e.g., timestamp, auto-increment)
- Low cardinality (e.g., country, status)
- Single-value dominated

Example:



javascript

// BAD: Monotonically increasing

```
sh.shardCollection("logs", { timestamp: 1 }) // All writes go to one shard
```

// BETTER: Hashed timestamp

```
sh.shardCollection("logs", { timestamp: "hashed" })
```

// BEST: Compound key

```
sh.shardCollection("logs", { user_id: 1, timestamp: 1 })
```

Chunk Management



javascript

// Check chunk distribution

```
db.chunks.find({ ns: "myDatabase.users" }).count()
```

// Split chunk

```
sh.splitAt("myDatabase.users", { user_id: 5000 })
```

// Move chunk

```
sh.moveChunk("myDatabase.users", { user_id: 1000 }, "shard2RS")
```

// Enable/disable balancer

```
sh.enableBalancing("myDatabase.users")
```

```
sh.disableBalancing("myDatabase.users")
```

// Check balancer status

```
sh.getBalancerState()
```

Targeted vs Broadcast Operations



javascript

// Targeted query (includes shard key)

```
db.users.find({ user_id: 12345 }) // Goes to specific shard
```

// Broadcast query (no shard key)

```
db.users.find({ email: "john@example.com" }) // Goes to all shards
```

// Targeted update

```
db.users.updateOne(  
  { user_id: 12345 }, // Includes shard key  
  { $set: { status: "active" } }  
)
```

// Broadcast update

```
db.users.updateMany(  
  { status: "inactive" }, // No shard key  
  { $set: { archived: true } }  
)
```

12. Security

Authentication

Enable Authentication



javascript

// Create admin user

```
use admin  
db.createUser({  
  user: "admin",  
  pwd: "securePassword123",  
  roles: ["root"]  
})
```



bash

Restart mongod with authentication

```
mongod --auth --port 27017 --dbpath /data/db
```

Connect with authentication

```
mongosh -u admin -p securePassword123 --authenticationDatabase admin
```

Create Database Users



javascript

// Create read-only user

```
use myDatabase
```

```
db.createUser({  
  user: "reader",  
  pwd: "password123",  
  roles: [{ role: "read", db: "myDatabase" }]  
})
```

// Create read-write user

```
db.createUser({  
  user: "writer",  
  pwd: "password123",  
  roles: [{ role: "readWrite", db: "myDatabase" }]  
})
```

// Custom roles

```
db.createUser({  
  user: "customUser",  
  pwd: "password123",  
  roles: [  
    { role: "read", db: "myDatabase" },  
    { role: "readWrite", db: "analytics" }  
  ]  
})
```

Built-in Roles



```
// Database roles
- read: Read data
- readWrite: Read and write data
- dbAdmin: Database administration
- dbOwner: Database owner (all privileges)

// Cluster roles
- clusterAdmin: Cluster administration
- clusterManager: Monitor and manage cluster
- clusterMonitor: Read-only access to monitoring tools

// Backup/Restore roles
- backup: Backup data
- restore: Restore data

// All-database roles
- readAnyDatabase: Read all databases
- readWriteAnyDatabase: Read/write all databases
- dbAdminAnyDatabase: Admin all databases
- userAdminAnyDatabase: Manage users in all databases

// Superuser role
- root: Full access to all resources
```

User Management



// List users

```
db.getUsers()
```

// View user

```
db.getUser("username")
```

// Update user password

```
db.changeUserPassword("username", "newPassword")
```

// Grant role to user

```
db.grantRolesToUser("username", [{ role: "readWrite", db: "myDatabase" }])
```

// Revoke role from user

```
db.revokeRolesFromUser("username", [{ role: "read", db: "myDatabase" }])
```

// Drop user

```
db.dropUser("username")
```

Network Security



bash

Bind to specific IP

```
mongod --bind_ip localhost,192.168.1.100
```

Enable TLS/SSL

```
mongod --tlsMode requireTLS \  
  --tlsCertificateKeyFile /path/to/mongod.pem \  
  --tlsCAFile /path/to/ca.pem
```

Encryption

Encryption at Rest



bash

```
mongod --enableEncryption \
  --encryptionKeyFile /path/to/keyfile
```

Encryption in Transit (TLS)



bash

```
mongod --tlsMode requireTLS \
  --tlsCertificateKeyFile /path/to/cert.pem
```

Field-Level Encryption



javascript

```
const { MongoClient, ClientEncryption } = require('mongodb')
```

```
// Configure encryption
```

```
const client = new MongoClient(uri, {
  autoEncryption: {
    keyVaultNamespace: 'encryption.__keyVault',
    kmsProviders: {
      local: {
        key: Buffer.from(localMasterKey, 'base64')
      }
    }
  }
})
```

```
// Insert with automatic encryption
```

```
await client.db('test').collection('users').insertOne({
  name: 'John',
  ssn: '123-45-6789' // Automatically encrypted
})
```

Auditing



bash

```
# Enable auditing
mongod --auditDestination file \
  --auditFormat JSON \
  --auditPath /var/log/mongodb/audit.json \
  --auditFilter '{ atype: { $in: ["authenticate", "createUser", "dropDatabase"] } }'
```



javascript

```
// Audit filter examples
// Log all authentication attempts
{ atype: "authenticate" }

// Log all write operations
{ atype: { $in: ["insert", "update", "delete"] } }

// Log operations on specific database
{ "param.ns": /^myDatabase\.\/ }
```

13. Performance Optimization

Query Optimization

1. Use Indexes



javascript

```
// Before (slow)
db.users.find({ email: "john@example.com" }) // COLLSCAN

// Create index
db.users.createIndex({ email: 1 })

// After (fast)
db.users.find({ email: "john@example.com" }) // IXSCAN
```

2. Use Projection



javascript

```
// Bad (retrieves all fields)
db.users.find({ age: { $gt: 25 } })

// Good (retrieves only needed fields)
db.users.find(
  { age: { $gt: 25 } },
  { name: 1, email: 1, _id: 0 }
)
```

3. Avoid Regex Without Index



javascript

```
// Very slow
db.users.find({ name: { $regex: /john/i } })

// Better: Use text index
db.users.createIndex({ name: "text" })
db.users.find({ $text: { $search: "john" } })

// Or anchor regex
db.users.find({ name: { $regex: /^john/i } }) // Can use index
```

4. Use Covered Queries



javascript

```
db.users.createIndex({ name: 1, age: 1 })

// Covered query (no document access needed)
db.users.find(
  { name: "John" },
  { name: 1, age: 1, _id: 0 }
)
```

5. Limit Results



javascript

```
// Bad  
db.users.find() // Returns all  
  
// Good  
db.users.find().limit(100)
```

6. Use \$in Efficiently



javascript

```
// Bad (multiple queries)  
db.users.find({ id: 1 })  
db.users.find({ id: 2 })  
db.users.find({ id: 3 })  
  
// Good (single query)  
db.users.find({ id: { $in: [1, 2, 3] } })  
  
// But limit size of $in array (< 1000 elements)
```

Schema Design for Performance

1. Embed for Read Performance



javascript

// One query to get everything

```
{
  _id: 1,
  name: "John",
  address: {
    street: "123 Main St",
    city: "NYC"
  },
  orders: [
    { id: 101, total: 50 },
    { id: 102, total: 75 }
  ]
}
```

2. Reference for Write Performance



javascript

// Users collection

```
{ _id: 1, name: "John" }
```

// Orders collection (many writes don't affect user doc)

```
{ _id: 101, user_id: 1, total: 50 }
```

```
{ _id: 102, user_id: 1, total: 75 }
```

3. Use Appropriate Data Types



javascript

// Bad

```
{ age: "25" } // String
```

// Good

```
{ age: 25 } // Number
```

// Bad

```
{ created_at: "2024-01-15" } // String
```