

# SQL Complete Study Material - Zero to Hero

---

## Chapter 1: Introduction to Databases and SQL

### What is a Database?

A database is a structured collection of data stored electronically. Think of it as a digital filing system where information is organized in tables.

**Real-world example:** A school database might contain:

- Student information (names, IDs, grades)
- Course information (course names, credits, professors)
- Enrollment records (which students are in which courses)

### What is SQL?

SQL (Structured Query Language) is the standard language for:

- **Querying** data (asking questions)
- **Inserting** new data
- **Updating** existing data
- **Deleting** unwanted data
- **Creating** database structures

### Database Terminology

**Table:** A collection of related data organized in rows and columns

employees table:

id	name	position	salary
1	Jhon	Developer	70000
2	Sarah	Manager	85000

**Row (Record):** A single entry in a table (like John's complete information) **Column (Field):** A specific attribute (like "name" or "salary") **Primary Key:** A unique identifier for each row (like employee ID)

---

## Chapter 2: Basic SELECT Statements

### The SELECT Statement Structure

```
SELECT column_name(s)
```

```
FROM table_name;
```

### Examples with Explanations

#### Example 1: Select all columns

```
SELECT * FROM employees;
```

- \* means "all columns"
- This returns every piece of information about every employee

#### Example 2: Select specific columns

```
SELECT first_name, salary FROM employees;
```

- Only shows names and salaries
- Useful when you don't need all information

#### Example 3: Select with aliases (renaming columns)

```
SELECT first_name AS "Employee Name", salary AS "Monthly Pay"
```

```
FROM employees;
```

- AS creates a temporary name for display
- Makes output more readable

### Practice Exercises - Chapter 2

1. Write a query to show all information about customers
  2. Show only customer names and cities
  3. Display product names with a more readable column header
  4. Select employee emails and hire dates
-

# Chapter 3: Filtering Data with WHERE

## The WHERE Clause

Used to filter rows based on conditions.

### Basic Syntax:

```
SELECT column_name(s)
```

```
FROM table_name
```

```
WHERE condition;
```

## Comparison Operators

- `=` Equal to
- `!=` or `<>` Not equal to
- `>` Greater than
- `<` Less than
- `>=` Greater than or equal
- `<=` Less than or equal

## Examples with Explanations

### Example 1: Exact match

```
SELECT * FROM employees WHERE department = 'Sales';
```

**Explanation:** Shows all employees who work in the Sales department

### Example 2: Numeric comparison

```
SELECT name, salary FROM employees WHERE salary > 50000;
```

**Explanation:** Shows names and salaries of employees earning more than \$50,000

### Example 3: Text patterns with LIKE

```
SELECT * FROM customers WHERE first_name LIKE 'J%';
```

**Explanation:**

- `LIKE` is used for pattern matching
- `%` means "any characters"
- `J%` means "starts with J"
- `%son` would mean "ends with son"
- `%mid%` would mean "contains mid"

## Logical Operators

### **AND - Both conditions must be true**

SELECT \* FROM employees

WHERE department = 'IT' AND salary > 60000;

### **OR - At least one condition must be true**

SELECT \* FROM customers

WHERE country = 'USA' OR country = 'Canada';

### **NOT - Opposite of the condition**

SELECT \* FROM orders WHERE NOT status = 'Cancelled';

### **IN - Match any value in a list**

SELECT \* FROM customers WHERE country IN ('USA', 'UK', 'Canada');

### **BETWEEN - Within a range**

SELECT \* FROM products WHERE price BETWEEN 10 AND 100;

## Practice Exercises - Chapter 3

1. Find all employees hired after 2020
  2. Show customers from either 'USA' or 'UK'
  3. Display products with price less than \$50
  4. Find employees whose email contains 'gmail'
  5. Show orders with amount between \$100 and \$500
-

# Chapter 4: Sorting and Limiting Results

## ORDER BY Clause

Used to sort results in ascending or descending order.

### Syntax:

```
SELECT column_name(s)

FROM table_name

ORDER BY column_name ASC|DESC;
```

## Examples

### Example 1: Sort by salary (highest first)

```
SELECT name, salary FROM employees

ORDER BY salary DESC;
```

### Example 2: Sort by multiple columns

```
SELECT * FROM employees

ORDER BY department ASC, salary DESC;
```

**Explanation:** First sorts by department A-Z, then within each department, sorts by salary highest to lowest

## LIMIT Clause

Used to restrict the number of rows returned.

### Example: Top 5 highest-paid employees

```
SELECT name, salary FROM employees

ORDER BY salary DESC

LIMIT 5;
```

## Practice Exercises - Chapter 4

1. Show all products sorted by price (cheapest first)
2. Display top 3 most recent orders
3. List customers alphabetically by last name
4. Show 10 most expensive products

---

## Chapter 5: Aggregate Functions

### What are Aggregate Functions?

Functions that perform calculations on multiple rows and return a single result.

### Common Aggregate Functions

#### **COUNT()** - Counts rows

```
SELECT COUNT(*) FROM employees;
```

-- Result: Total number of employees

```
SELECT COUNT(DISTINCT department) FROM employees;
```

-- Result: Number of different departments

#### **SUM()** - Adds up values

```
SELECT SUM(salary) FROM employees;
```

-- Result: Total of all salaries

```
SELECT SUM(quantity) FROM order_items;
```

-- Result: Total items sold

#### **AVG()** - Calculates average

```
SELECT AVG(salary) FROM employees;
```

-- Result: Average salary

```
SELECT AVG(price) FROM products;
```

-- Result: Average product price

### **MAX() and MIN() - Find highest and lowest values**

```
SELECT MAX(salary), MIN(salary) FROM employees;
```

-- Result: Highest and lowest salaries

```
SELECT MAX(order_date) FROM orders;
```

-- Result: Most recent order date

### **Combining Aggregates with WHERE**

```
SELECT AVG(salary) FROM employees
```

```
WHERE department = 'IT';
```

-- Average salary in IT department only

### **Practice Exercises - Chapter 5**

1. Count total number of customers
  2. Find the most expensive product
  3. Calculate average order amount
  4. Sum all product stock quantities
  5. Find earliest employee hire date
-

## Chapter 6: GROUP BY and HAVING

### GROUP BY Clause

Groups rows with the same values and allows aggregate functions on each group.

#### Example: Average salary by department

```
SELECT department, AVG(salary) as avg_salary
```

```
FROM employees
```

```
GROUP BY department;
```

#### Result:

department	avg_salary
sales	65000
it	78000
HR	55000

### More GROUP BY Examples

#### Example: Count employees by department

```
SELECT department, COUNT(*) as employee_count
```

```
FROM employees
```

```
GROUP BY department;
```

#### Example: Total sales by customer

```
SELECT customer_id, SUM(total_amount) as total_spent
```

```
FROM orders
```

```
GROUP BY customer_id;
```

### HAVING Clause

Used to filter groups (like WHERE but for groups).

#### Example: Departments with more than 5 employees



SELECT department, COUNT(\*) as employee\_count

FROM employees

GROUP BY department

HAVING COUNT(\*) > 5;

**WHERE vs HAVING:**

- **WHERE** filters individual rows before grouping
- **HAVING** filters groups after grouping

**Practice Exercises - Chapter 6**

1. Count orders by status
  2. Average product price by category
  3. Find customers who spent more than \$1000 total
  4. Show departments with average salary > \$60000
-

# Chapter 7: JOINS - Combining Tables

## Why JOINS?

Real-world data is spread across multiple related tables. JOINS combine data from these tables.

## Types of JOINS

### INNER JOIN

Returns rows that have matching values in both tables.

**Example: Employee names with department names**

```
SELECT e.first_name, e.last_name, d.department_name  
FROM employees e  
INNER JOIN departments d ON e.department_id = d.department_id;
```

#### Explanation:

- **e** and **d** are table aliases (shortcuts)
- **ON** specifies how tables are related
- Only employees with valid departments are shown

### LEFT JOIN

Returns all rows from the left table, plus matched rows from the right table.

**Example: All employees with their departments (even if no department)**

```
SELECT e.first_name, d.department_name  
FROM employees e  
LEFT JOIN departments d ON e.department_id = d.department_id;
```

### RIGHT JOIN

Returns all rows from the right table, plus matched rows from the left table.

**Example: All departments with their employees (even empty departments)**

```
SELECT e.first_name, d.department_name  
FROM employees e  
RIGHT JOIN departments d ON e.department_id = d.department_id;
```

## Complex JOIN Example

**Orders with customer and product information:**

```
SELECT  
    c.first_name + ' ' + c.last_name as customer_name,  
    p.product_name,  
    oi.quantity,  
    o.order_date  
FROM orders o  
INNER JOIN customers c ON o.customer_id = c.customer_id  
INNER JOIN order_items oi ON o.order_id = oi.order_id  
INNER JOIN products p ON oi.product_id = p.product_id;
```

## Practice Exercises - Chapter 7

1. Show employee names with their manager names
  2. List all orders with customer information
  3. Display products with their order quantities
  4. Find customers who haven't placed any orders
-

# Chapter 8: Subqueries

## What is a Subquery?

A query inside another query. Used when you need the result of one query to help with another.

## Simple Subquery Example

**Find employees earning more than average:**

```
SELECT first_name, last_name, salary  
  
FROM employees  
  
WHERE salary > (SELECT AVG(salary) FROM employees);
```

### Step-by-step explanation:

1. Inner query calculates: `SELECT AVG(salary) FROM employees` → Result: 70000
2. Outer query becomes: `WHERE salary > 70000`
3. Shows employees with salary > 70000

## Subquery with IN

**Find customers who have placed orders:**

```
SELECT first_name, last_name  
  
FROM customers  
  
WHERE customer_id IN (SELECT customer_id FROM orders);
```

## Practice Exercises - Chapter 8

1. Find products more expensive than average price
  2. Show customers who placed orders in 2023
  3. Find employees in departments with more than 3 people
-

## Chapter 9: Advanced Features

### CASE Statements

Used for conditional logic (like if-else).

**Example: Categorize employees by salary**

```
SELECT
    first_name,
    salary,
    CASE
        WHEN salary < 50000 THEN 'Low'
        WHEN salary < 80000 THEN 'Medium'
        ELSE 'High'
    END as salary_category
FROM employees;
```

### UNION

Combines results from multiple SELECT statements.

**Example: All email addresses from customers and employees**

```
SELECT email FROM customers
UNION
SELECT email FROM employees;
```

### Window Functions (Advanced)

Perform calculations across related rows.

**Example: Rank employees by salary within department**

```
SELECT
```

first\_name,

department\_id,

salary,

RANK() OVER (PARTITION BY department\_id ORDER BY salary DESC) as salary\_rank

FROM employees;

---

## Chapter 10: Data Modification

### INSERT - Adding New Data

-- Insert single row

```
INSERT INTO customers (first_name, last_name, email)
VALUES ('John', 'Doe', 'john@email.com');
```

-- Insert multiple rows

```
INSERT INTO products (product_name, price, category) VALUES
('Laptop', 999.99, 'Electronics'),
('Mouse', 19.99, 'Electronics');
```

### UPDATE - Modifying Existing Data

-- Update single record

```
UPDATE employees
SET salary = 75000
WHERE employee_id = 101;
```

-- Update multiple records

```
UPDATE products
SET price = price * 1.1
WHERE category = 'Electronics';
```

### DELETE - Removing Data

-- Delete specific records

```
DELETE FROM orders
WHERE status = 'Cancelled';
```

-- Delete all records (be careful!)

```
DELETE FROM temp_table;
```

---

## Chapter 11: Database Design and Normalization

### What is Database Design?

Database design is the process of organizing data efficiently and logically. Poor design leads to data redundancy, inconsistency, and performance issues.

### Database Normalization

Normalization is the process of organizing data to reduce redundancy and improve data integrity.

#### First Normal Form (1NF)

**Rule:** Each column should contain atomic (indivisible) values, and each row should be unique.

#### Bad Example (Not 1NF):

students table:

Id	Name	subject
1	Jhon	Maths, Science
2	Jane	English, History

#### Good Example (1NF):

students table:

Id	Name
1	Jhon
2	Jane

student\_subjects table:



Id	Student_id	Subject
1	1	Math
2	1	Science
3	2	English
4	2	History

## Second Normal Form (2NF)

**Rule:** Must be in 1NF, and all non-key columns must depend on the entire primary key.

**Example:** In an order\_details table with composite primary key (order\_id, product\_id), customer\_name shouldn't be there because it only depends on order\_id, not the full key.

## Third Normal Form (3NF)

**Rule:** Must be in 2NF, and no non-key column should depend on another non-key column.

### Bad Example:

employees table:

Id	Name	Dept	Dept_location	Salary
1	Jhon	Sales	New York	50000
2	Jane	Sales	New York	55000

*Problem: dept\_location depends on dept, not on employee\_id*

### Good Example (3NF):

employees table:

ID	Name	dept_id	Salary
1	Jhon	1	50000
2	Jane	1	55000

departments table:

Dept_id	Dept_location
1	New York
2	London

## Entity Relationship (ER) Diagrams

Visual representation of database structure showing:

- **Entities** (tables) - rectangles
- **Attributes** (columns) - ovals
- **Relationships** - diamonds

### Types of Relationships

1. **One-to-One (1:1)**: Each record in table A relates to exactly one record in table B
    - Example: Employee ↔ Employee\_Details
  2. **One-to-Many (1:M)**: One record in table A can relate to many records in table B
    - Example: Department → Employees
  3. **Many-to-Many (M:M)**: Records in both tables can relate to multiple records in the other
    - Example: Students ↔ Courses (requires junction table)
-

# Chapter 12: Constraints and Data Integrity

## Primary Key Constraints

```
CREATE TABLE employees (  
    employee_id INT PRIMARY KEY,  
    name VARCHAR(50) NOT NULL  
);
```

## Foreign Key Constraints

```
CREATE TABLE orders (  
    order_id INT PRIMARY KEY,  
    customer_id INT,  
    FOREIGN KEY (customer_id) REFERENCES customers(customer_id)  
);
```

## Check Constraints

```
CREATE TABLE products (  
    product_id INT PRIMARY KEY,  
    price DECIMAL(10,2) CHECK (price > 0),  
    stock_quantity INT CHECK (stock_quantity >= 0)  
);
```

## Unique Constraints

```
CREATE TABLE users (  
    user_id INT PRIMARY KEY,  
    email VARCHAR(255) UNIQUE,
```

```
username VARCHAR(50) UNIQUE  
);
```

## NOT NULL Constraints

```
CREATE TABLE customers (  
    customer_id INT PRIMARY KEY,  
    first_name VARCHAR(50) NOT NULL,  
    last_name VARCHAR(50) NOT NULL,  
    email VARCHAR(255) NOT NULL UNIQUE  
);
```

---

## Chapter 13: Indexes and Performance

### What are Indexes?

Indexes are database objects that improve query performance by creating shortcuts to data.

**Analogy:** Like an index in a book - instead of reading every page to find "SQL", you look in the index and jump directly to page 245.

### Creating Indexes

-- Create index on frequently searched column

```
CREATE INDEX idx_customer_email ON customers(email);
```

-- Composite index for multiple columns

```
CREATE INDEX idx_order_date_status ON orders(order_date, status);
```

-- Unique index

```
CREATE UNIQUE INDEX idx_product_code ON products(product_code);
```

## When to Use Indexes

### Good for:

- Columns frequently used in WHERE clauses
- Columns used in JOIN conditions
- Columns used in ORDER BY

### Avoid indexing:

- Small tables (< 1000 rows)
- Columns that change frequently
- Tables with heavy INSERT/UPDATE activity

## Query Performance Tips

-- Good: Uses index

```
SELECT * FROM customers WHERE customer_id = 123;
```

-- Bad: Function prevents index usage

```
SELECT * FROM customers WHERE UPPER(email) = 'JOHN@EMAIL.COM';
```

-- Good: Index-friendly

```
SELECT * FROM customers WHERE email = 'john@email.com';
```

---

# Chapter 14: Views and Stored Procedures

## Views

A view is a virtual table based on a query. It doesn't store data but shows results dynamically.

-- Create a view

```
CREATE VIEW employee_summary AS
```

```
SELECT
```

```
    e.first_name + ' ' + e.last_name as full_name,
```

```
    d.department_name,
```

```
    e.salary
```

```
FROM employees e
```

```
JOIN departments d ON e.department_id = d.department_id;
```

-- Use the view like a table

```
SELECT * FROM employee_summary WHERE salary > 60000;
```

### Benefits of Views:

- Simplify complex queries
- Security (hide sensitive columns)
- Consistent interface for applications

## Stored Procedures

Reusable SQL code blocks that can accept parameters.

-- Create stored procedure

```
CREATE PROCEDURE GetEmployeesByDepartment(IN dept_name VARCHAR(50))
```

```
BEGIN
```

```
    SELECT e.first_name, e.last_name, e.salary
```

```
    FROM employees e
```

```
JOIN departments d ON e.department_id = d.department_id  
WHERE d.department_name = dept_name;  
END;
```

```
-- Execute stored procedure
```

```
CALL GetEmployeesByDepartment('Sales');
```

---

# Chapter 15: Transactions and ACID Properties

## What are Transactions?

A transaction is a sequence of SQL operations treated as a single unit of work.

-- Start transaction

```
BEGIN TRANSACTION;
```

-- Multiple operations

```
UPDATE accounts SET balance = balance - 100 WHERE account_id = 1;
```

```
UPDATE accounts SET balance = balance + 100 WHERE account_id = 2;
```

-- Commit if everything succeeds

```
COMMIT;
```

-- Or rollback if something fails

```
-- ROLLBACK;
```

## ACID Properties

**A - Atomicity:** All operations succeed or all fail **C - Consistency:** Database remains in valid state **I - Isolation:** Concurrent transactions don't interfere **D - Durability:** Committed changes are permanent

---



## Chapter 16: Common Table Expressions (CTEs)

### What are CTEs?

Temporary named result sets that exist only during query execution.

-- Simple CTE

```
WITH sales_summary AS (  
    SELECT  
        customer_id,  
        SUM(total_amount) as total_sales  
    FROM orders  
    GROUP BY customer_id  
)  
SELECT  
    c.first_name,  
    c.last_name,  
    s.total_sales  
FROM customers c  
JOIN sales_summary s ON c.customer_id = s.customer_id  
WHERE s.total_sales > 1000;
```

### Recursive CTEs

Used for hierarchical data like organizational charts.

-- Find all employees under a manager

```
WITH RECURSIVE employee_hierarchy AS (  
    -- Base case: Start with the manager  
    SELECT employee_id, first_name, manager_id, 0 as level  
    FROM employees
```

WHERE employee\_id = 101

UNION ALL

-- Recursive case: Find direct reports

SELECT e.employee\_id, e.first\_name, e.manager\_id, eh.level + 1

FROM employees e

JOIN employee\_hierarchy eh ON e.manager\_id = eh.employee\_id

)

SELECT \* FROM employee\_hierarchy;

---

# Chapter 17: JSON and Modern SQL Features

## Working with JSON Data

Many modern databases support JSON data types.

-- PostgreSQL JSON examples

```
CREATE TABLE products (
```

```
    id INT PRIMARY KEY,
```

```
    name VARCHAR(100),
```

```
    attributes JSON
```

```
);
```

-- Insert JSON data

```
INSERT INTO products VALUES
```

```
(1, 'Laptop', '{"brand": "Apple", "ram": "16GB", "storage": "512GB"}');
```

-- Query JSON data

```
SELECT name, attributes->>'brand' as brand
```

```
FROM products
```

```
WHERE attributes->>'ram' = '16GB';
```

## Window Functions (Advanced Analytics)

Perform calculations across related rows without grouping.

-- Running total of sales

```
SELECT
```

```
    order_date,
```

```
    total_amount,
```

```
    SUM(total_amount) OVER (ORDER BY order_date) as running_total
```

FROM orders;

-- Rank within groups

SELECT

department\_id,

first\_name,

salary,

RANK() OVER (PARTITION BY department\_id ORDER BY salary DESC) as dept\_rank

FROM employees;

---

## Important Concepts Summary

### Data Types to Know

- **Numeric:** INT, DECIMAL(10,2), FLOAT
- **Text:** VARCHAR(50), TEXT, CHAR(10)
- **Date/Time:** DATE, DATETIME, TIMESTAMP
- **Boolean:** BOOLEAN (TRUE/FALSE)
- **Modern:** JSON, XML, ARRAY

### SQL Standards and Dialects

- **ANSI SQL:** Standard that all databases follow
- **PostgreSQL:** Advanced features, strong JSON support
- **MySQL:** Popular for web applications
- **SQL Server:** Microsoft's enterprise solution
- **Oracle:** Enterprise-grade with advanced features

### Security Best Practices

1. **Use parameterized queries** to prevent SQL injection
2. **Grant minimum necessary permissions** to users
3. **Encrypt sensitive data** at rest and in transit
4. **Regular backups** and disaster recovery plans
5. **Audit database access** and changes